



CSC380: Principles of Data Science

**Introduction to Machine Learning /
Basics of Predictive Modeling and
Classification**

Kyoungseok Jang

Announcement

- Delay: Due to the urgent circumstances of the TA in charge, we haven't finished the grading of the Midterm and HW4. We will have it done by next Tuesday. HW5 is out now (due : 3/24)
- Midterm Curving
 - I am thinking about $\sqrt{100 \times (\text{Your Score})}$ as the curved score.
 - E.g.) If your score is 50, your curved score is slightly over 70.

Because of the TA's circumstances,
One problem is not graded yet.



I expect around 50% after full grading...

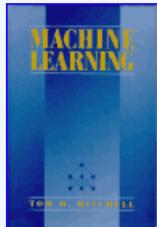
- Self-Withdrawal deadline: 3/28

What is machine learning?

3

- Tom Mitchell established Machine Learning Department at CMU (2006).

Machine Learning, Tom Mitchell, McGraw Hill, 1997. “**through experience**”



Machine Learning is the study of computer algorithms that improve automatically through experience. Applications range from datamining programs that discover general rules in large data sets, to information filtering systems that automatically learn users' interests.

This book provides a single source introduction to the field. It is written for advanced undergraduate and graduate students, and for developers and researchers in the field. No prior background in artificial intelligence or statistics is assumed.

- A bit outdated with recent trends, but still has interesting discussion (and easy to read).
- A subfield of Artificial Intelligence – you want to perform nontrivial, smart tasks. The difference from the traditional AI is “how” you build a computer program to do it.

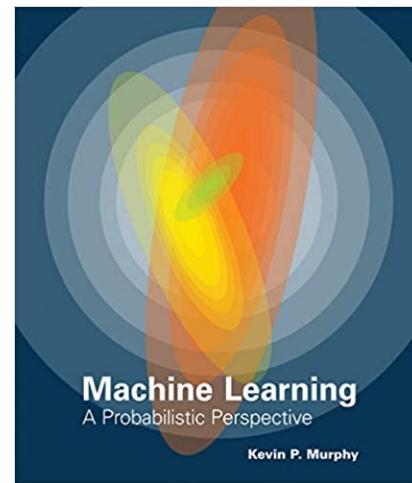
Textbooks

4

We will use a more recent textbook for readings

*Takes a **probabilistic approach** to
machine learning*

*Consistent with the goals of data
science in this class*



Murphy, K. "Machine Learning: A Probabilistic Perspective." MIT press, 2012

[\(UA Library\)](#)

AI Task 1: Image classification

5

- Predefined categories: $\mathcal{C} = \{\text{cat, dog, lion, ...}\}$
- Given an image, classify it as one of the categories $c \in \mathcal{C}$ with the highest accuracy.
- Use: sorting/searching images by category, medical imaging, object identification, traffic control, categorizing types of stars/events in the Universe (images taken from large surveying telescopes)



AI Task 2: Recommender systems

6

- Predict how user would rate a movie
- Use: For each user, pick an unwatched movie with high predicted ratings. (Youtube, Netflix, Amazon, etc.)
- Idea: compute user-user similarity or movie-movie similarity, then compute a weighted average.



	User 1	User 2	User 3
Movie 1	1	2	1
Movie 2	?	3	1
Movie 3	2	5	2
Movie 4	4	?	5
Movie 5	?	4	5

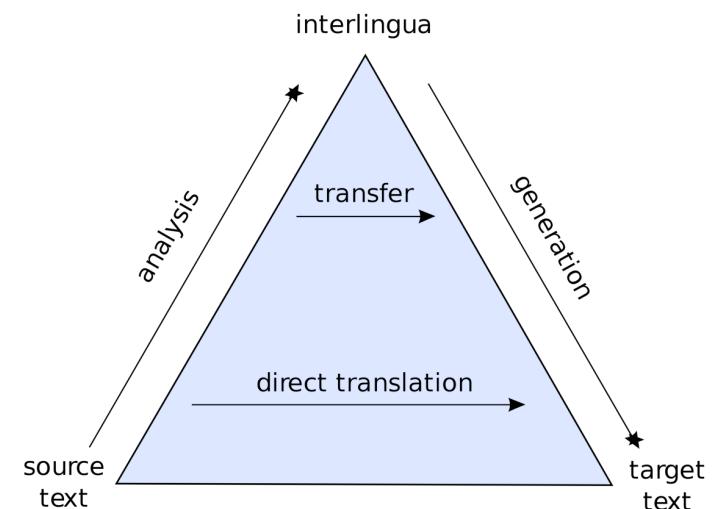
“collaborative filtering”

AI Task 3: Machine translation

7

- No need to explain how useful it is.
- **Task:** 1) Transform a sentence to the interlingual language (analysis) and 2) create a sentence with another language with the same meaning, with appropriate grammar structure (generation).

A screenshot of a machine translation application. At the top, there are language selection dropdowns: '언어 감지' (Language detection), '한국어' (Korean), '영어' (English, highlighted in blue), '중국어(간체)' (Simplified Chinese), and '영어' (English). Below this is a double-headed arrow icon. To the right are more dropdowns: '영어' (English), '한국어' (Korean), and '중국어(간체)' (Simplified Chinese). The main area shows a translation pair: 'Hello world!' in English and '안녕 세상!' in Korean. Below the Korean text is its phonetic transcription: 'annyeong sesang!'. At the bottom left are audio playback icons, and at the bottom center is a progress bar showing '12 / 5,000' with a dropdown arrow next to it. On the far right are additional icons.



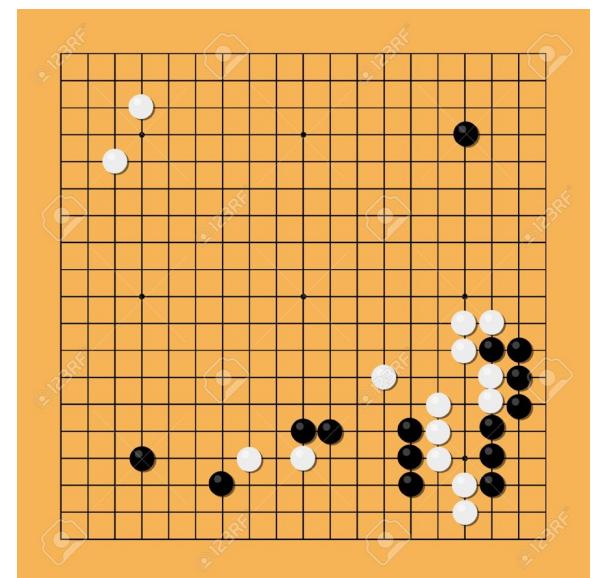
AI Task 4: Board game

8

- Predict win probability of a move in a given game state (e.g., AlphaGo)
- Traditionally considered as a “very smart” task to perform.

Q: how will it be useful for us, though?

- Use: From the AI Go player, you can do practice play or even learn from it.
 - Now it's a major trend in the field of Go
- Potential use: Board game (e.g., Catan) design, better AI
- Deeply related to robot AI and autonomous driving
 - Predict the future of your move



Traditional AI vs Machine Learning (ML)



9

- **Traditional AI:** you encode the knowledge (e.g., logic statements/rules), and the *machine* executes it.
 - e.g., if there is feather-like texture with two eyes and a beak, classify it as a bird.
 - Advancements in automated ‘**inference**’ like “if $a \rightarrow b$ and $b \rightarrow c$, then $a \rightarrow c$ ”. \Rightarrow ‘expert system’
- **ML:** Given a set of input and output pairs (e.g., animal picture + label), and train a **function** (a set of logical statements / a neural network) that maps the input to the output accurately.
 - As the “big data” era comes, data is abundant \Rightarrow turns out, better than systems based on hand-coded domain knowledge!
 - “statistical” approach // data-driven approach

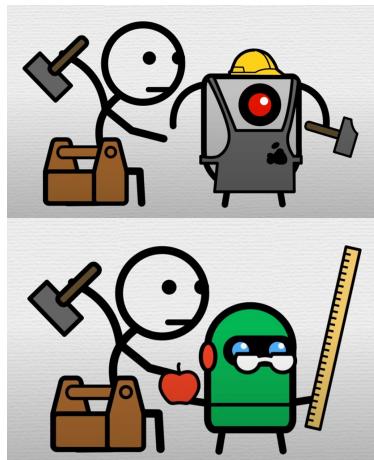
“Every time I fire a linguist, the performance of the speech recognizer goes up.”

– 1988, Frederick Jelinek, a Czech-American researcher in information theory & speech recognition.

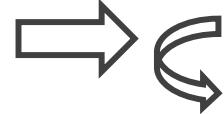
Traditional AI vs Machine Learning (ML)

10

- Traditional AI – watchmaker
 - You encode your knowledge (springs and parts) directly
 - You understand why those parts are necessary.
- Machine Learning (ML) – one example (from <https://www.youtube.com/watch?v=R9OHn5ZF4Uo>)

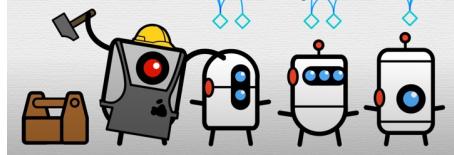


Auto-updater



Auto-tester

Automatically improve through
Interaction and experience



It is possible that
you might not know
why your function is
designed in that way



Overview of ML Methods

11

Supervised Learning

- Provide *training* data consisting of input-output pairs and learn mapping
- E.g., Spam prediction, object detection or image classification, machine translation, etc.

Unsupervised learning

- No predefined categories. Finds patterns in the data without the help of labels (outputs)
- E.g., clustering, dimensionality reduction, target tracking, image segmentation, etc.

Reinforcement learning **We won't cover this**

- The environment interacts with your action, transferring you to different states.
- E.g., autonomous driving, robot AI, recommendation system

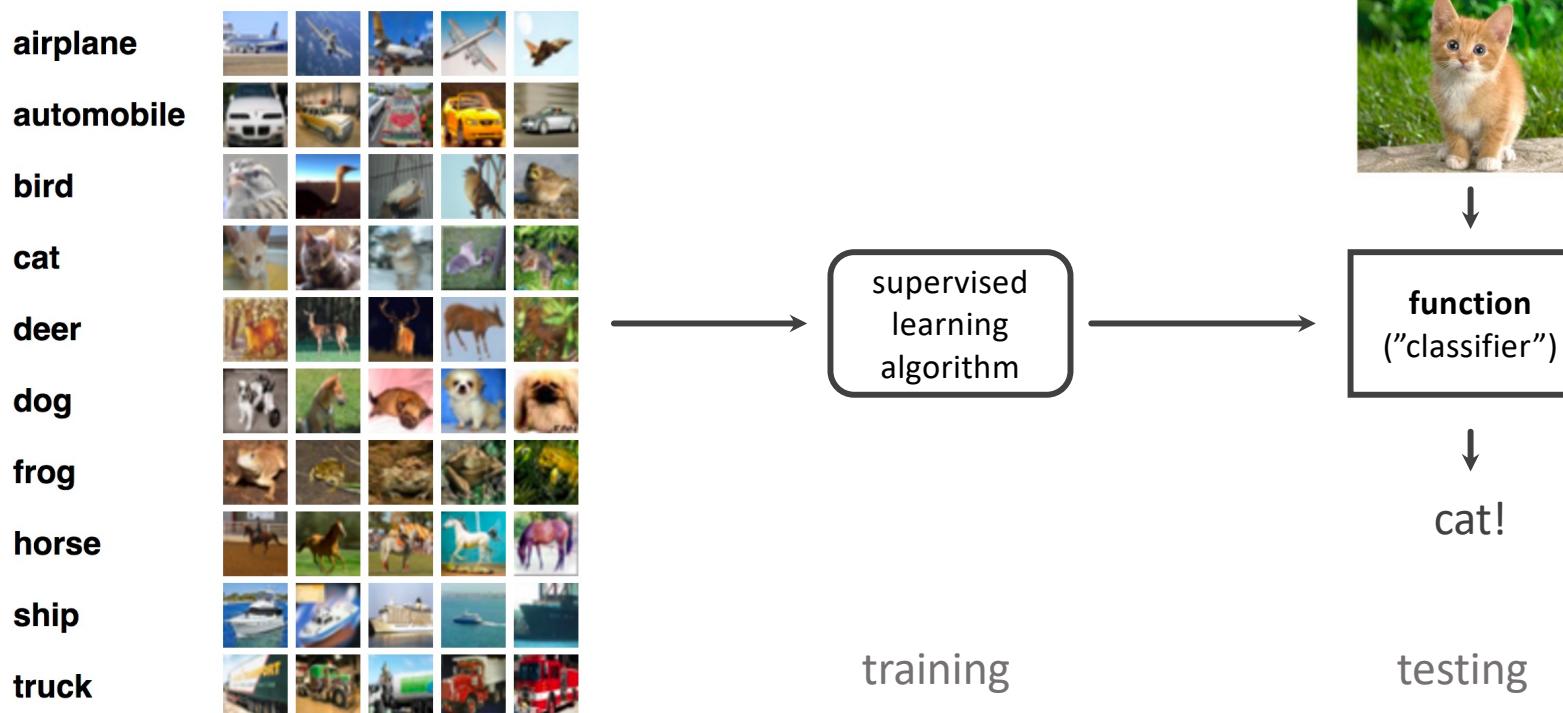
Supervised Learning

Basic setting: Supervised learning

example = data point
labeled = categorized

13

- Train data: dataset comprised of *labeled examples*: a pair of (input, label)



Example function 1: Decision tree

14

```
Task: predict the 5-star rating of a movie by a user
```

```
If age >= 60 then
    if genre = western then
        return 4.3
    else if release date > 1998 then
        return 2.5
    else ...
    ...
end if

else if age < 60 then
...
end if
```

training:

- determine the shape of the tree
- which condition to have at each node
- what to output from each leaf node

Example function 2: Linear

Task: Image classification

Let x be a set of pixel values of a picture (30x30) =>
900-dimensional vector x .

If $0.124 \cdot x_1 - 2.5 \cdot x_2 + \dots + 2.31 \cdot x_{900} - 2.12 \geq 0$ then
return cat

else

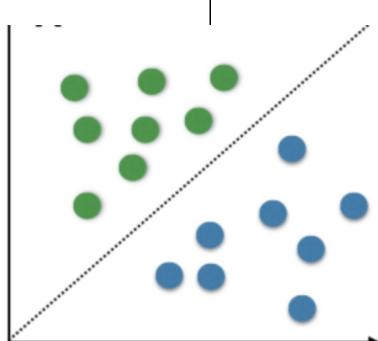
return dog

end

“linear combination”/“inner product”

training:

- determine the coefficients & threshold

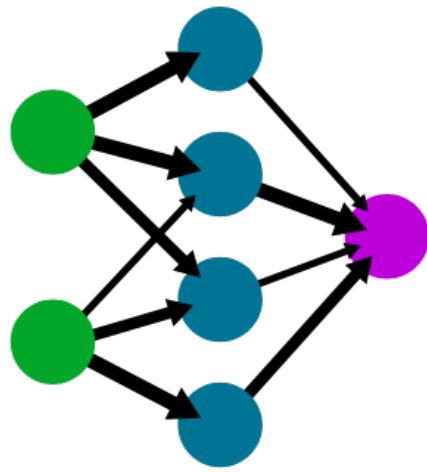


E.g., in 2d space, it induces a linear decision boundary:
 $0.5 \cdot x_1 - 2.5 \cdot x_2 > 4.3$

Example function 3: Nonlinear

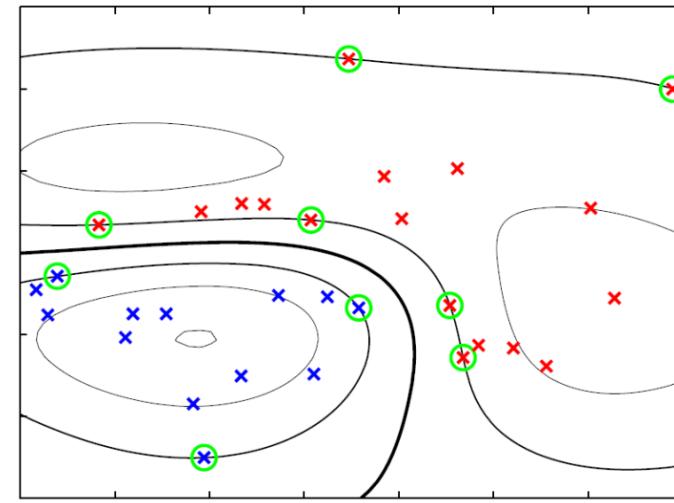
16

Neural network



(stacked **linear** models with nonlinear **activation functions**)

Support Vector Machine



(**linear** in the induced feature space)

Example: Naïve Bayes Classifier

Training Data:

Person	height (feet)	weight (lbs)	foot size(inches)
male	6	180	12
male	5.92 (5'11")	190	11
male	5.58 (5'7")	170	12
male	5.92 (5'11")	165	10
female	5	100	6
female	5.5 (5'6")	150	8
female	5.42 (5'5")	130	7
female	5.75 (5'9")	150	9



Task: Observe feature vector $x = (x_1, \dots, x_n)$ and predict class label $y \in \{1, \dots, C\}$

Model: Treat features as *conditionally independent*, given class label:

$$p(x, y) = p(y) \prod_{i=1}^n p(x_i|y)$$

Doesn't capture correlation among features, but is easier to learn.

Classification: Bayesian model so classify by posterior,

$$p(y = c | x) = \frac{p(C = k)p(x|y = c)}{p(x)}$$

Supervised learning: Types of prediction problems ¹⁸

Binary classification: Choose between 2 classes

- Given an email, is it spam or not? (or the probability of it being a spam)

Multi-class classification: more than 2 categories.

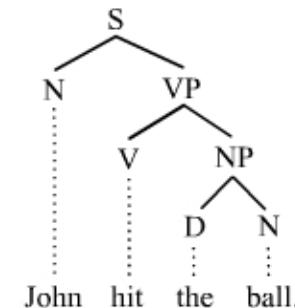
- Image classification with 1000 categories. (cat, dog, airplane, car, computer, ...)

Regression: the label is real-valued (e.g., price)

- Say I am going to visit Italy next month. Given the price trends in the past, what would be the price given (the # of days before the departure, day of week)?
- Predict the stocks/bitcoin price in the future

Structured output prediction: more than just a number

- Given a sentence, what is its grammatical parse tree?

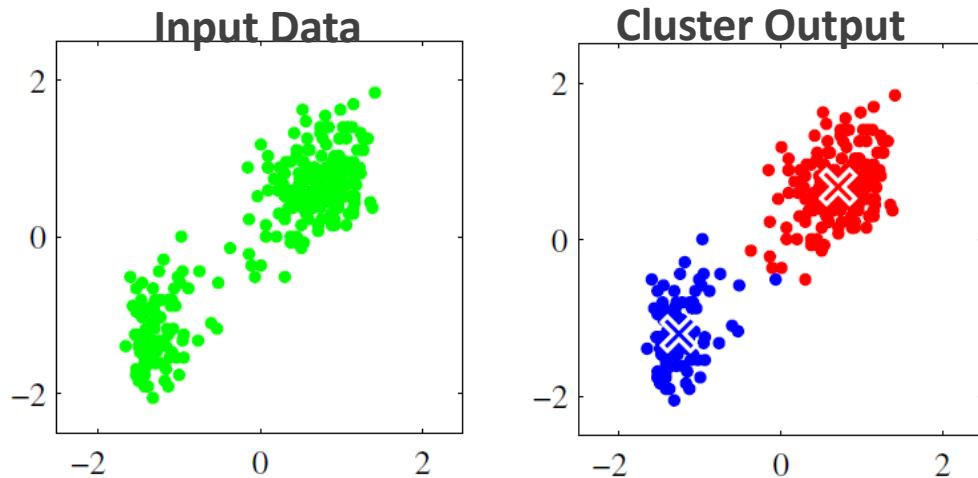


Unsupervised Learning

Example: Clustering

20

Identify groups (clusters) of similar data



Useful for interpreting large datasets

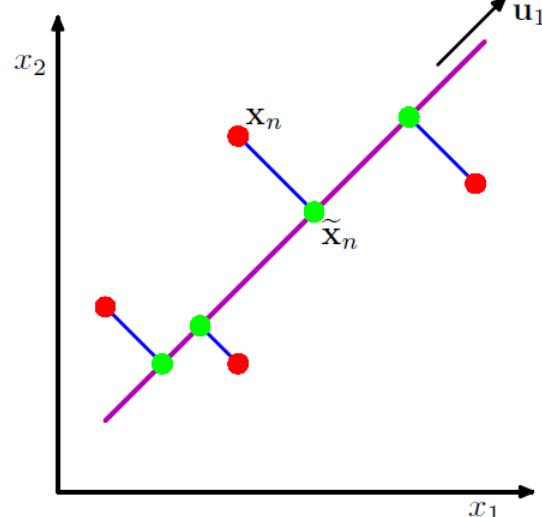
Clusters are assigned arbitrary labels (e.g. 1, 2, ..., K).
=> afterwards, you may look at the data and name each group.

Common clustering algorithms: K-means, Expectation Maximization (EM)

Example: Principal Component Analysis (PCA)

21

Reduce dimension of high-dimensional data using linear projection



Identify directions of **maximum variation** in the data by computing **eigenvectors**

Easier explanation: Identify important directions

Linear projection onto K-dimensional subspace spanned by top K eigenvalues

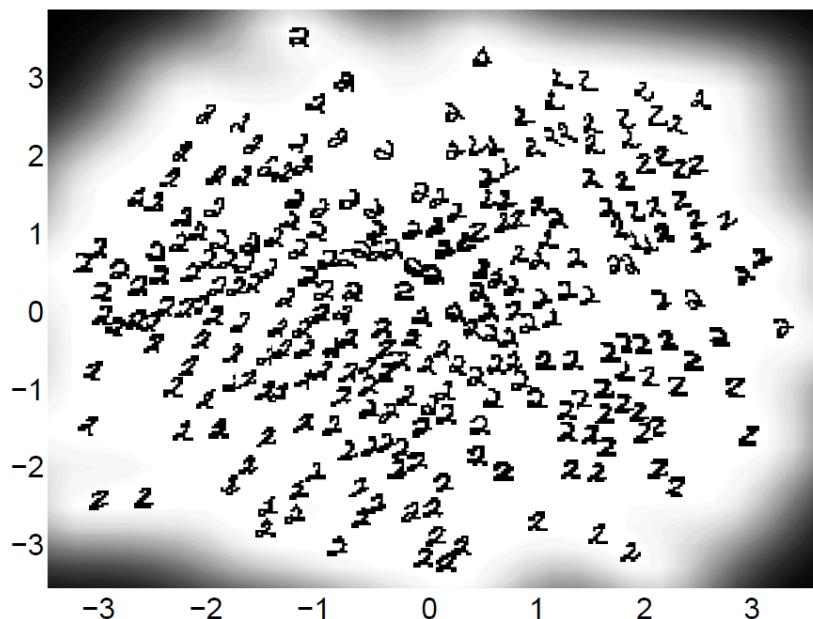
Can be used for visualization (project to 2D) or for compressing images.

Source: Bishop, C. PRML

Example: Principal Component Analysis (PCA)

22

Reduce dimension of high-dimensional data using linear projection

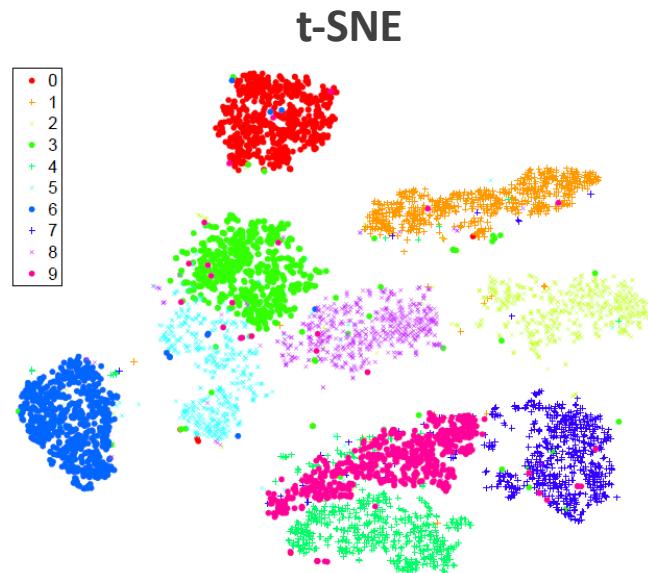


Source: Lawrence, N. (2005)

Example for modeling / visualizing
handwritten digits

Each digit is a black/white image with
28x28 pixels (784 dimensions)
projected down to 2D

Example: Nonlinear Dimensionality Reduction 23



Nonlinear reduction can (potentially) amplify clustering properties

t-Distributed Stochastic Neighbor Embedding (t-SNE) Models similarity between data as a t distribution and strives to find projection that preserves similarity.

Example: Generative models



We won't cover this

- AI image generators
- It is hard to define how ‘good’ the generated image is.
 - How can we explain the ‘painting style’ to computers? Mostly impossible... → Unsupervised!

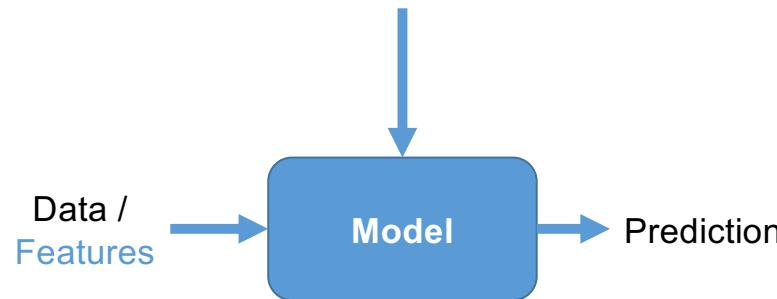


- **Supervised Learning** - Training data consist of inputs and outputs
 - Classification, regression, translation, ...
- **Unsupervised Learning** – Training data only contain inputs
 - Clustering, dimensionality reduction, segmentation, ...
- **Linear** models generate output as a *linear combination* of inputs,
 - E.g. $y = w_1x_1 + w_2x_2 + \dots + w_dx_d$
 - PCA, linear regression, etc.
- **Nonlinear** models fit an arbitrary nonlinear function to map inputs-outputs
 - Neural networks, support vector machine, nonlinear dimensionality reduction

Supervised Learning

works with labeled data

Labels / Outputs



Unsupervised Learning

works with unlabeled data



ML models distinguished by a number of factors

- Number of parameters needed (parametric / nonparametric)
- Whether they model uncertainty (probabilistic / nonprobabilistic)
- Do they model the data generation process? (generative / discriminative)

CSC380: Principles of Data Science

Basics of Predictive Modeling and Classification 1: Decision Tree

Kyoungseok Jang

Decision Trees

figures/examples from “A Course in Machine Learning” by Hal Daume III <http://ciml.info/>

Majority Vote Classifier

29

The most basic classifier you can think of.

How to train:

- Given: A (train) dataset with m data points $\{(x^{(i)}, y^{(i)})\}_{i=1}^m$ with C classes.
- Compute the most common class c^* in the dataset.
- Output a classifier $f(x) = c^*$.

$$c^* = \arg \max_{c \in \{1, \dots, C\}} \sum_{i=1}^m \mathbf{I}\{y^{(i)} = c\}$$

Example:

Data: $m=10$

$x^{(i)}$: images of cats and dogs

$y^{(i)}$: label (cat/dog)

Suppose that there are 6 dogs and 4 cats.

After ‘training’, your classifier always outputs ‘dog’, even without looking at the input.

Stupid enough classifier! Always try to beat this classifier.

Often, state-of-the-art ML algorithms perform barely better than the majority vote classifier..

⇒ happens when there is no association between features and labels in the dataset

Train set accuracy

30

- Suppose the ML algorithm has trained a function f using the dataset $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^m$ where $x^{(i)}$ is input and $y^{(i)}$ is label.
- Train set accuracy:

$$\widehat{acc}(f) := \frac{1}{m} \sum_{i=1}^m \mathbf{I}\{f(x^{(i)}) = y^{(i)}\}$$

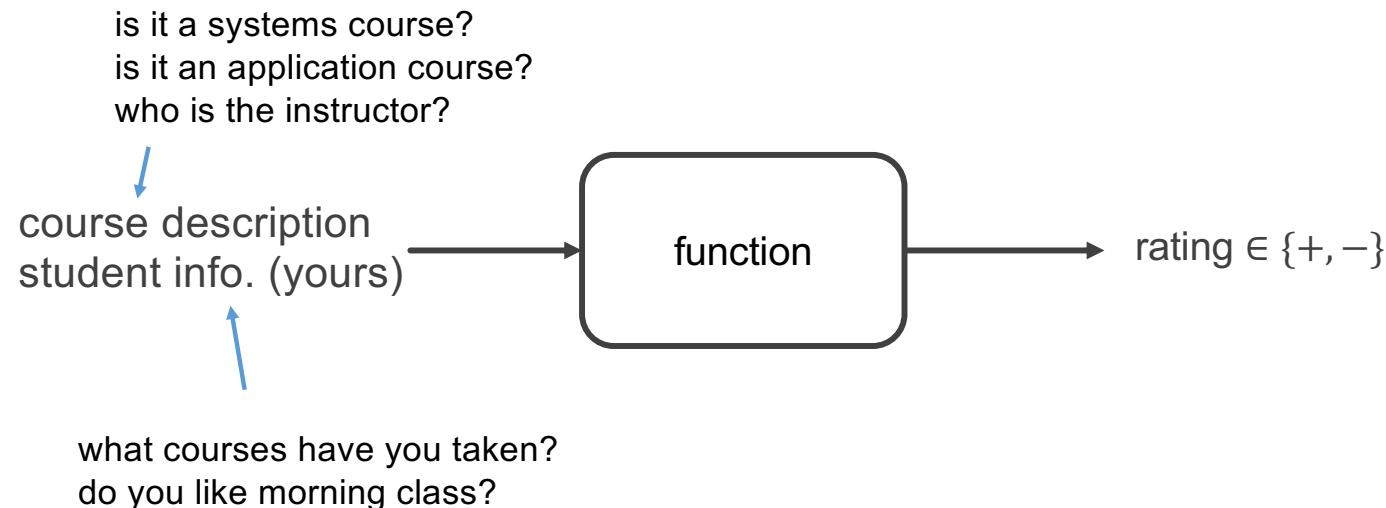
It is the number of times
the function got the
answer right divided by m.

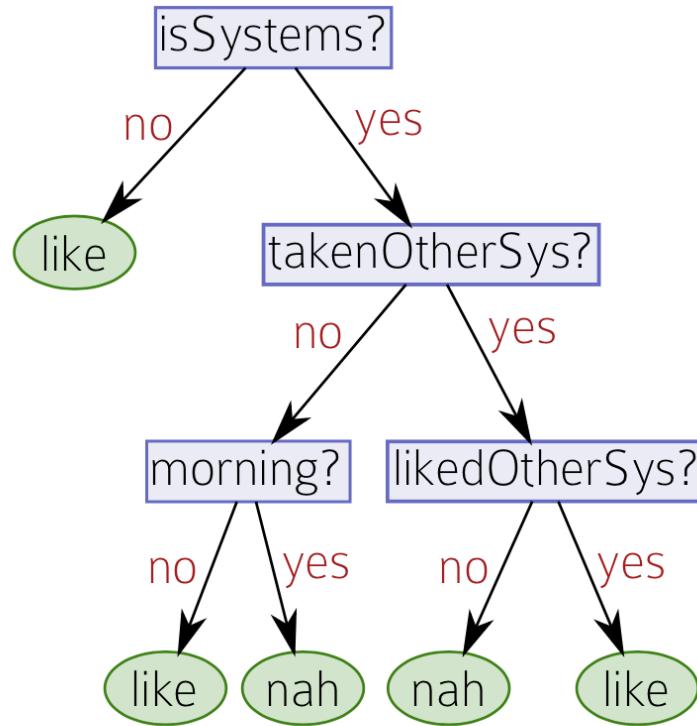
- Q: We have 100 data points (images) with 5 cats, 80 dogs, and 15 lions. What is the train set accuracy of the majority vote classifier?

.80

Decision tree (example: course recommendation) 32

- Build software: recommend a set of courses for you
 - More precisely, given a course, predict its rating





Wouldn't it be nice to construct such a tree automatically by a computer algorithm?

Wouldn't it be nice if it accurately predicts?

You can, if you have data!

Rating	Easy?	AI?	Sys?	Thy?	Morning?
consider it to be 'like'	+2	y	y	n	y
	+2	y	y	n	y
	+2	n	y	n	n
	+2	n	n	n	y
	+2	n	y	y	n
	+1	y	y	n	n
	+1	y	y	n	y
	+1	n	y	n	y
	0	n	n	n	n
	0	y	n	n	y
consider it to be 'dislike'	0	n	y	n	y
	-1	y	y	y	n
	-1	n	n	y	y
	-1	n	n	y	n
	-1	y	n	y	n
	-2	n	n	y	y
	-2	n	y	y	n
	-2	y	n	y	n
	-2	y	n	y	y

HasTakenPrereqs (=: Prereq)
 ↓
 HasTakenACourseFromTheSameLecturer (=: Lecturer)
 ↓
 HasLabs

For example, this table is data D.
 Each row is a course you've rated.
 $x^{(i)}$ is a sequence of 5 yes/no ($d=5$) for i -th course.
 $y^{(i)}$ is the sign of the rating for i -th course.

Define the data $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^m$

$\in \{y, n\}^d$
 $\in \{+, -\}$

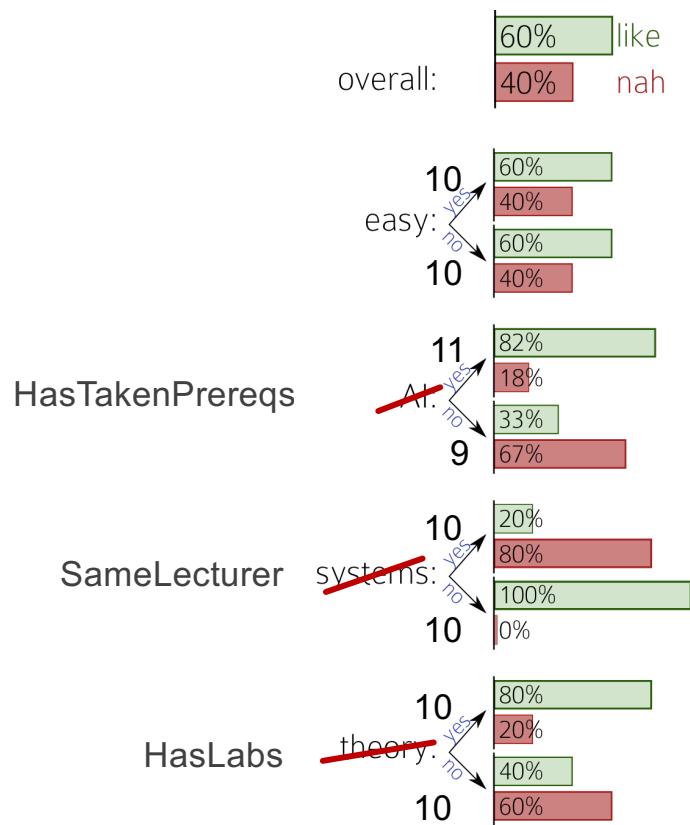
Each dimension of $x^{(i)}$ is called a **feature**.
 $x^{(i)}$ is called a **feature vector**.

- Main principle: Find a tree that has a high train set accuracy

$$\widehat{acc}(f) = \frac{1}{m} \sum_{i=1}^m \mathbf{I}\{f(x^{(i)}) = y^{(i)}\}$$

- This is essentially the main principle governing pretty much all the machine learning algorithms!
 - “[Empirical risk minimization](#)” principle
(empirical risk := 1 – train_accuracy)

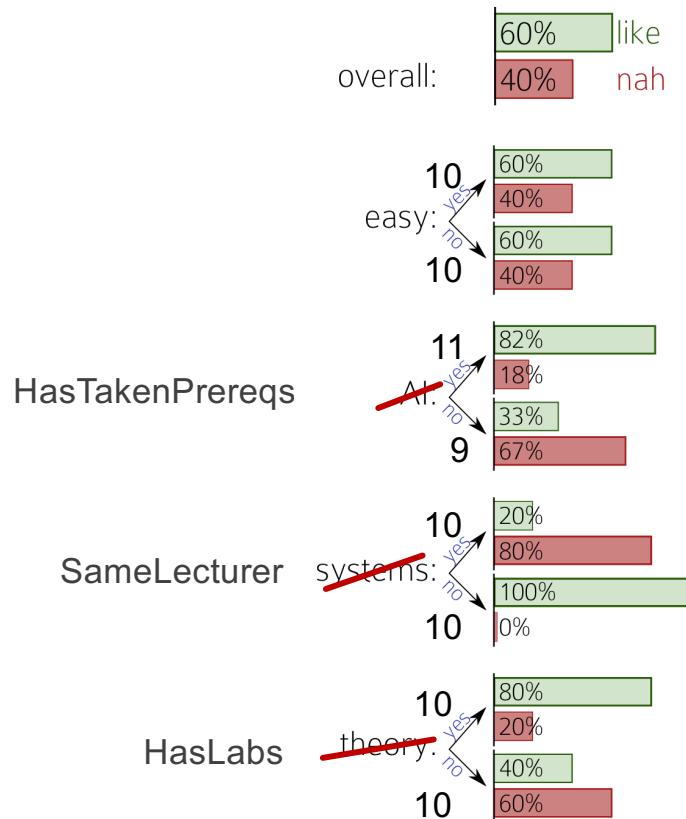
How to construct a tree



Rating	Easy?	Prereqs	Lecturer	HasLabs	Morning?
		AI?	Sys?	Theory?	
+2	y	y	n	y	n
+2	y	y	n	y	n
+2	n	y	n	n	n
+2	n	n	n	y	n
+2	n	y	y	n	y
+1	y	y	n	n	n
+1	y	y	n	y	n
+1	n	y	n	y	n
0	n	n	n	n	y
0	y	n	n	y	y
0	n	y	n	y	n
0	y	y	y	y	y
-1	y	y	y	n	y
-1	n	n	y	y	n
-1	n	n	y	n	y
-1	y	n	y	n	y
-2	n	n	y	y	n
-2	n	y	y	n	y
-2	y	n	y	n	n
-2	y	n	y	n	y

How to construct a tree

37

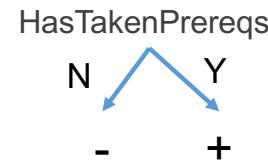


Baseline: majority vote classifier

Q: What is the train set accuracy? 0.60

Major
(+)

Suppose we place the node **HasTakenPrereqs** at the root.
Set the prediction at each leaf node as the majority vote.

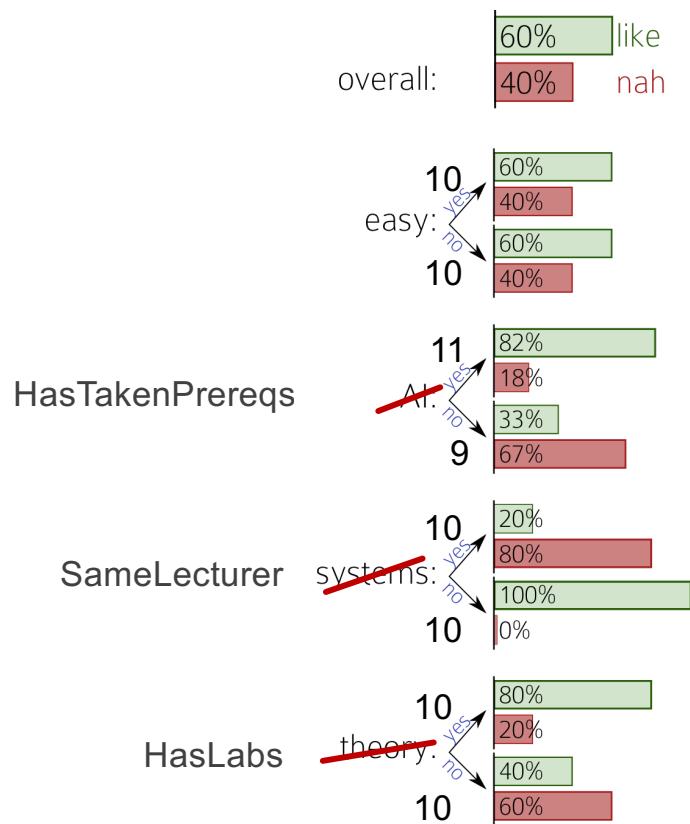


What is the train set accuracy now?

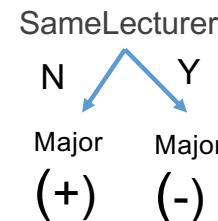
$$\frac{9}{20} \cdot \frac{6}{9} + \frac{11}{20} \cdot \frac{9}{11} = \frac{15}{20} = 0.75 \quad \text{improved!}$$

How to construct a tree

38



Suppose placing the node SameLecturer at the root.



What is the train set accuracy now?

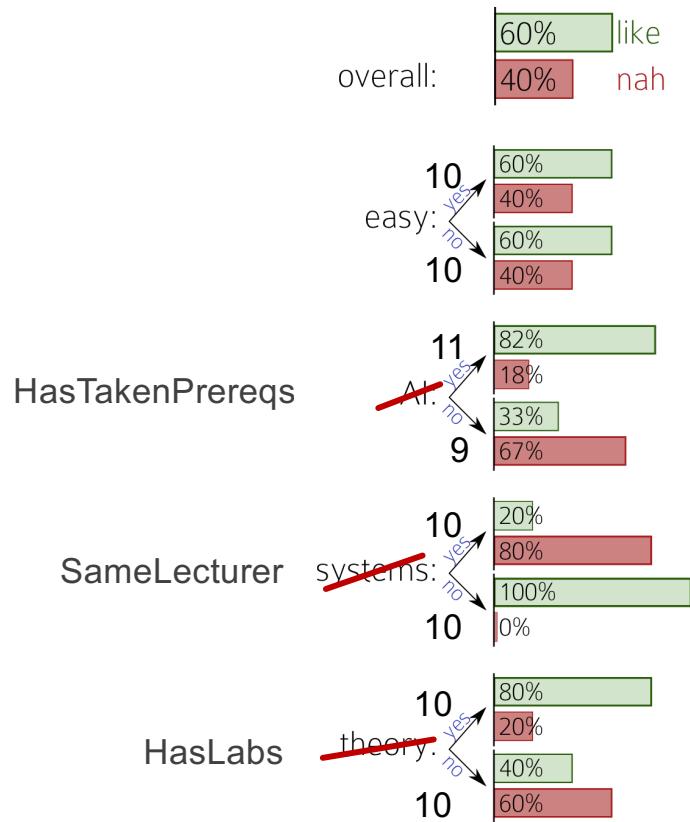
$$\frac{10}{20} \cdot \frac{10}{10} + \frac{10}{20} \cdot \frac{8}{10} = \frac{18}{20} = 0.9 \quad \text{even better!}$$

What would you do to build a depth-1 tree?

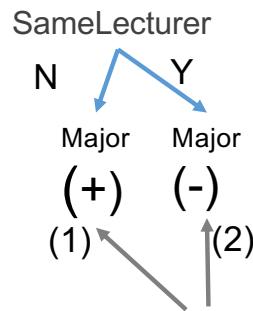
try out each feature and choose the one that leads to the largest accuracy!

How to construct a tree

39



What about depth 2?

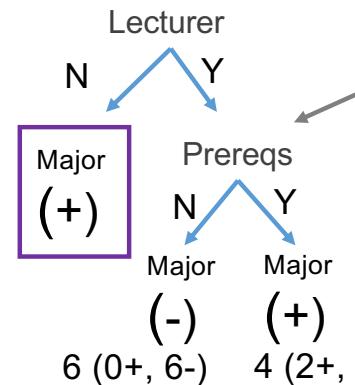
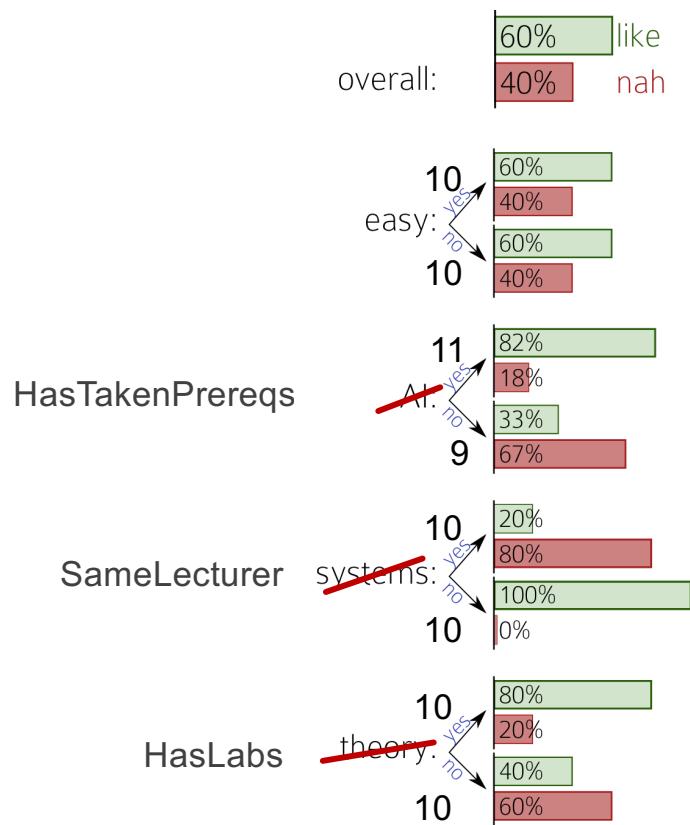


Which nodes to put at each leaf node?

Focus on (2). Try placing HasTakenPrereqs

How to construct a tree

40



Q: How many training data points fall here? 10

Q: How many training data points arrive at these two leaves? How many for each label?

Q: what prediction should we use for each leaf?

Q: What is the train set accuracy, conditioning on
SameLecturer=Y?

'local' train set accuracy

$$\frac{6}{10} \cdot \frac{6}{6} + \frac{4}{10} \cdot \frac{2}{4} = \frac{8}{10}$$

Try all the other nodes and pick the one with the largest acc.!

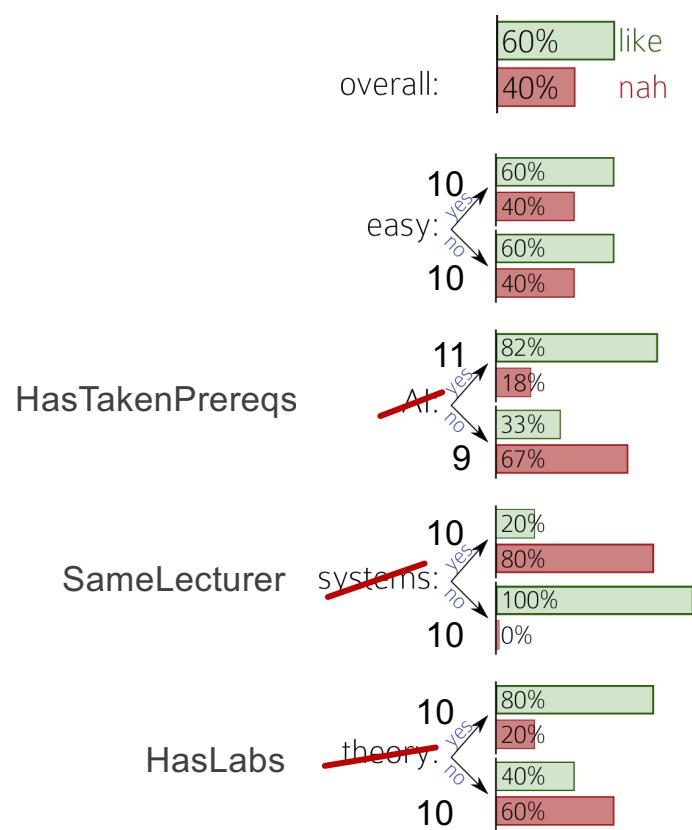
Then, repeat the same for **SameLecturer=N** branch!

=> but this has 1 local train set acc. So leave it be!

Move onto expanding nodes at depth 2!

How to construct a tree

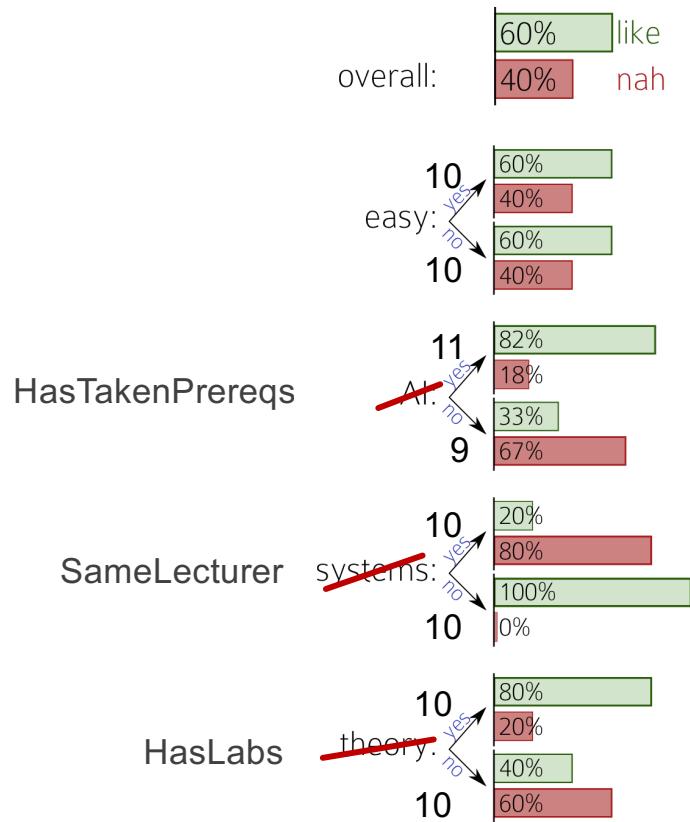
41



Rating	Easy?	Prereqs	Lecturer	HasLabs	
		AI?	Sys?	Theory?	Morning?
+2	y	y	n	y	n
+2	y	y	n	y	n
+2	n	y	n	n	n
+2	n	n	n	y	n
+2	n	y y		n	y
+1	y	y	n	n	n
+1	y	y	n	y	n
+1	n	y	n	y	n
0	n	n	n	n	y
0	y	n	n	y	y
0	n	y	n	y	n
0	y	y y		y	y
-1	y	y y		n	y
-1	n	n y		y	n
-1	n	n y		n	y
-1	y	n y		n	y
-2	n	n y		y	n
-2	n	y y		n	y
-2	y	n y		n	n
-2	y	n y		n	y

How to construct a tree

42



Overall idea:

1. Set the root node as a leaf node.
2. Grab a leaf node for which its 'local' train accuracy is not 1.
3. Find a feature that maximizes the 'local' train accuracy and replace the leaf node with a node with that feature; add leaf nodes and set their predictions by majority vote.
4. Repeat 2-3.

Algorithm 1 DECISIONTREETRAIN(*data*, *remaining features*)

```

1: guess  $\leftarrow$  most frequent answer in data           // default answer for this data
2: if the labels in data are unambiguous then           <= i.e., all data points have the same label
3:   return LEAF(guess)                                // base case: no need to split further
4: else if remaining features is empty then
5:   return LEAF(guess)                                // base case: cannot split further
6: else                                                 // we need to query more features
7:   for all f  $\in$  remaining features do           <= there is no point in adding a feature
8:     NO  $\leftarrow$  the subset of data on which f=no
9:     YES  $\leftarrow$  the subset of data on which f=yes
10:    score[f]  $\leftarrow$  ( # of majority vote answers in NO
11:      + # of majority vote answers in YES ) / size(data)           <= answer = label

12:   end for
13:   f  $\leftarrow$  the feature with maximal score(f)
14:   NO  $\leftarrow$  the subset of data on which f=no
15:   YES  $\leftarrow$  the subset of data on which f=yes
16:   left  $\leftarrow$  DECISIONTREETRAIN(NO, remaining features \ {f})
17:   right  $\leftarrow$  DECISIONTREETRAIN(YES, remaining features \ {f})
18:   return NODE(f, left, right)
19: end if

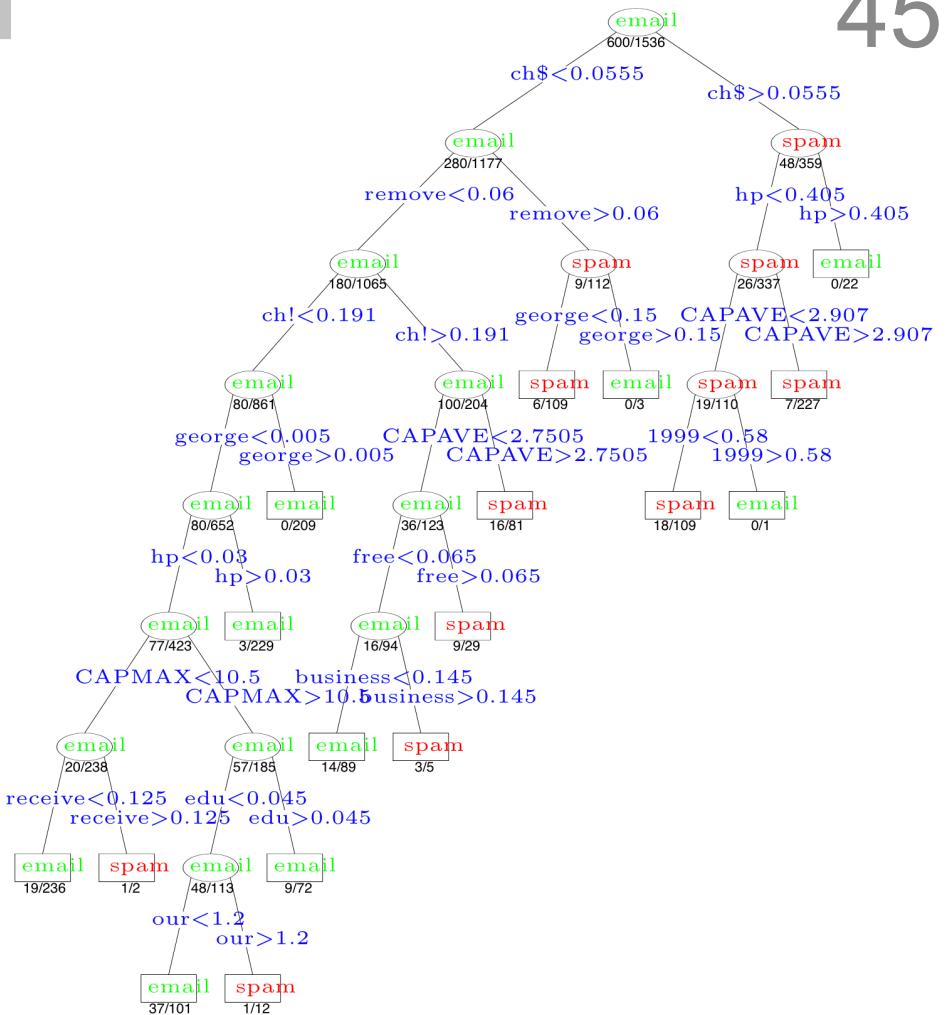
```

Algorithm 2 DECISIONTREETEST(*tree, test point*)

```
1: if tree is of the form LEAF(guess) then
2:   return guess
3: else if tree is of the form NODE(f, left, right) then
4:   if f = no in test point then
5:     return DECISIONTREETEST(left, test point)
6:   else
7:     return DECISIONTREETEST(right, test point)
8:   end if
9: end if
```

Example: spam filtering I

- ▶ Spam dataset
- ▶ 4601 email messages, about 39% are spam
- ▶ Classify message by spam and not-spam
- ▶ 57 features
 - ▶ 48 are of the form “percentage of email words that is (WORD)”
 - ▶ 6 are of the form “percentage of email characters is (CHAR)”
 - ▶ 3 other features (e.g., “longest sequence of all-caps”)
- ▶ Final tree after pruning has 17 leaves, 9.3% test error rate



Error := $1 - \text{accuracy}$.

Suppose we have trained a function \hat{f} on $D = \{(x^{(i)}, y^{(i)})\}_{i=1}^m$ using a supervised learning algorithm.

- Train error: Evaluate on D.

$$\widehat{\text{err}}_D(\textcolor{brown}{f}) := \frac{1}{|D|} \sum_{(x,y) \in D} \mathbf{I}\{\textcolor{brown}{f}(x) \neq y\}$$

- Test error: Evaluate on $D' = \{(x^{(i)}, y^{(i)})\}_{i=1}^{m'}$ not used for training.
 - It can be possible that our function just 'memorized' the training data and doesn't do well in real life. (overfitting)

Q: Choose one:

- (1) train error \geq test error (2) train error \approx test error (3) train error \leq test error

Background: Workflow of Training a Classifier

48

Standard practice:

- Given a data set D , split it into train set D_{train} and D_{test}
 - large data: 90-10 ratio
 - medium data: 80-20 ratio
 - small data: 70-30 ratio
- Train on D_{train} and evaluate error rate on D_{test} . You trust that D_{test} will be the performance when you deploy the trained classifier.
(these are guidelines only)

Discussion: What would be reasonable logics behind such a trust?

CSC380: Principles of Data Science

**Basics of Predictive Modeling and
Classification 2**

**Decision Trees / k-Nearest Neighborhood
Kyoungseok Jang**

- Final curving: you will recover 66% of the score you lost.
 - E.g.) If your original score was 40, your curved score will be $40 + (100 - 40) \times \frac{2}{3} = 80$
 - A bit more beneficial for the students who didn't have basic knowledge in probability and statistics.
 - New average: 83.3
- Final exam: I will spend a lecture for the final review, and I will try to 'describe' the problems more explicitly.
 - I will reuse several midterm problems with a bit of variation.

- Regrade request
 - Problem 7(4): We decided to give everyone the score. Please check your answer, and if your answer was 'False', please send us the regrade request.
 - For the student who used the back side of the paper for your answer, please let us know.

- Some students asked me about prerequisites.
 - Especially about the dimensionality reduction part
 - We will not evaluate you based on those prerequisites
 - We will not ask you like, how to calculate the eigenvectors or eigenvalues on your exam, or in your final project.
 - We will teach you the basic knowledge to understand. (E.g. inner product)
 - We will introduce you to some scipy functions for eigenvector computations.
 - (I am not sure whether we can cover the dimensionality reduction part)

I guess I will spend this lecture
for the final review...

Apr 25	Clustering 2
Apr 27	Dimensionality reduction 1
May 2	Dimensionality reduction 2
May 8	Final Exam (3:30-5:30pm)

- There was a gap between our scheduled progress and our current progress
 - For example, we should have finished the ‘sampling bias’ part before the midterm.
 - Therefore, now HW5 was posted too early.
 - We will learn k-Nearest Neighborhood today and Naïve-Bayes Classifier on Thursday. (both of them were included in HW5...)
- Therefore, we will extend the due date for HW5 to Mar. 31st.

- Due to his personal circumstances, TA Saiful will no longer be in charge of this class.
- Temporarily, we will not be able to provide the following services.
 - His office hour: Wed
 - His piazza hour: Wed/ Thu/ Fri

- Decision Tree
 - Review
 - Variations - Different criterions
 - Different types of features / labels
 - Regression
 - Pruning
- K-Nearest Neighborhood
 - Main concepts
 - Feature scaling
 - Variations / Issues

Review: Decision Tree method

56

Algorithm 1 DECISIONTREETRAIN(*data*, *remaining features*)

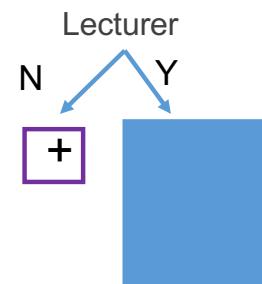
```
1: guess  $\leftarrow$  most frequent answer in data           // default answer for this data
2: if the labels in data are unambiguous then          <= i.e., all data points have the same label
3:   return LEAF(guess)                                // base case: no need to split further
4: else if remaining features is empty then
5:   return LEAF(guess)                                // base case: cannot split further
6: else                                                 // we need to query more features
7:   for all f  $\in$  remaining features do             <= there is no point in adding a feature
8:     NO  $\leftarrow$  the subset of data on which f=no        that appeared in its parent!
9:     YES  $\leftarrow$  the subset of data on which f=yes
10:    score[f]  $\leftarrow$  (# of majority vote answers in NO      <= score[f]='local' train set acc.
11:       + # of majority vote answers in YES ) / size(data)
12:   end for
13:   f  $\leftarrow$  the feature with maximal score(f)
14:   NO  $\leftarrow$  the subset of data on which f=no
15:   YES  $\leftarrow$  the subset of data on which f=yes
16:   left  $\leftarrow$  DECISIONTREETRAIN(NO, remaining features \ {f})
17:   right  $\leftarrow$  DECISIONTREETRAIN(YES, remaining features \ {f})
18:   return NODE(f, left, right)
19: end if
```

- Main question: How to calculate the ‘local’ train set accuracy? (*score*[*f*])

Review: Decision tree - How to construct a tree

57

Rating	Easy?	AI?	Sys?	Thy?	Morning?
+2	y	y	n	y	n
+2	y	y	n	y	n
+2	n	y	n	n	n
+2	n	n	n	y	n
+2	n	y	y	n	y
+1	y	y	n	n	n
+1	y	y	n	y	n
+1	n	y	n	y	n
0	n	n	n	n	y
0	y	n	n	y	y
0	n	y	n	y	n
0	y	y	y	y	y
-1	y	y	y	n	y
-1	n	n	y	y	n
-1	n	n	y	n	y
-1	y	n	y	n	y
-2	n	n	y	y	n
-2	n	y	y	n	y
-2	y	n	y	n	n
-2	y	n	y	n	y



Suppose that we already have this structure
And need to construct more on the RHS.

According to the algorithm, we need to calculate score[f],
which means the 'local' train set accuracy for each feature f.

Suppose that now our f is 'Prereqs'

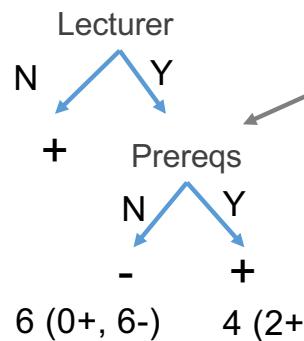
Main trick for calculating score[f]:

$$\begin{aligned} \widehat{acc}(f) &= \frac{1}{m} \sum_{i=1}^m \mathbf{I}\{f(x^{(i)}) = y^{(i)}\} \\ &= \frac{1}{m} (\sum_{i \in YES} \mathbf{I}\{f(x^{(i)}) = y^{(i)}\} + \sum_{i \in NO} \mathbf{I}\{f(x^{(i)}) = y^{(i)}\}) \end{aligned}$$

Review: Decision tree - How to construct a tree

58

Rating	Easy?	Prereqs	Lecturer		HasLabs	Morning?
			AI?	Sys?		
+2	y	y	n	y	n	
+2	y	y	n	y	n	
+2	n	y	n	n	n	
+2	n	n	n	y	n	
+2	n	y	y	n	y	
+1	y	y	n	n	n	
+1	y	y	n	y	n	
+1	n	y	n	y	n	
0	n	n	n	n	y	
0	y	n	n	y	y	
0	n	y	n	y	n	
0	y	y	y	y	y	
-1	y	y	y	n	y	
-1	n	n	y	y	n	
-1	n	n	y	n	y	
-1	y	n	y	n	y	
-2	n	n	y	y	n	
-2	n	y	y	n	y	
-2	y	n	y	n	n	
-2	y	n	y	n	y	



Q: How many training data points fall here? 10

Q: How many training data points arrive at these two leaves? How many for each label?

Q: what prediction should we use for each leaf?
A: Majority vote for each leaf

Q: How many samples will your current function outputs the 'correct' rating (sign) for each leaf?

Hint: Majority, since it is based on the majority vote 6 for left
2 for right

Q: What is the train set accuracy, conditioning on SameLecturer=Y?

'local' train set accuracy $\frac{1}{10}(6 + 2) = \frac{8}{10}$ or $(\frac{6}{10} \frac{6}{6} + \frac{4}{10} \frac{2}{4}) = \frac{8}{10}$

Sum of (fraction of sub group * fraction of correct answer in sub group)

Variations

59

Notions of uncertainty: binary case ($\mathcal{Y} = \{0, 1\}$)

Suppose in a set of examples $S \subseteq \mathcal{X} \times \{0, 1\}$, a p fraction are labeled as 1

What we did: Majority fraction = 1 – minority fraction

① Classification error: (red)

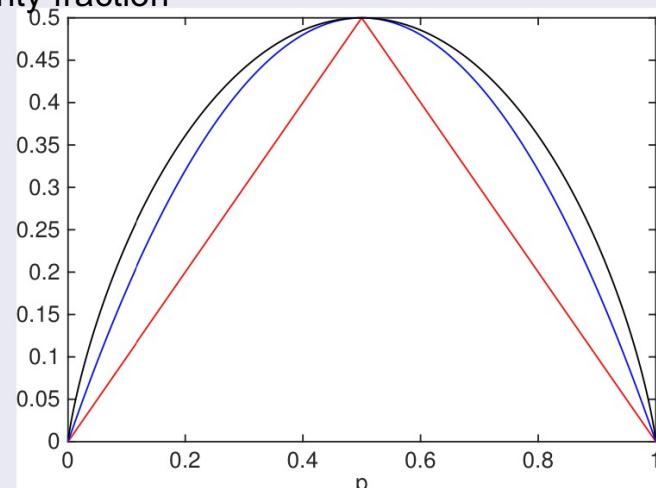
$$u(S) := \min\{p, 1 - p\}$$

② Gini index: (blue)

$$u(S) := 2p(1 - p)$$

③ Entropy: (black)

$$u(S) := p \log \frac{1}{p} + (1-p) \log \frac{1}{1-p}$$



Gini index and entropy (after some rescaling) are concave upper-bounds on classification error

Let q is the fraction of data points with feature=Y.

Modification:

Set score[f] as

$$q \cdot (-u(YES)) + (1 - q) \cdot (-u(NOT))$$

Set score[f] as $-u$ where u is one of these tree measure.

Using classification error is equivalent to using the accuracy.

For example,
in the previous example,
Lecturer=Y Prereqs=Y has
 $p=2/4$

Check: When u is the classification error, $q \cdot (-u(YES)) + (1 - q) \cdot (-u(NOT)) = (\text{score we knew}) - 1$

Decision tree – different types of features

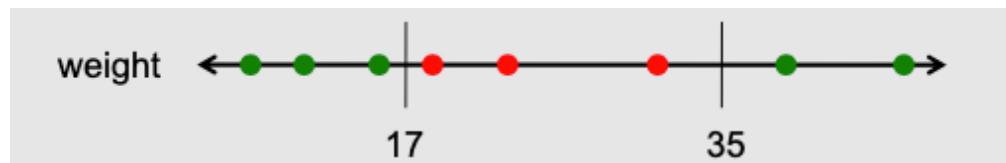
60

- Binary
- Categorical: values in $\{1, \dots, C\}$ e.g., occupation, blood type
 - Option 1: Instead of 2 children, have C children.
 - Option 2: Derive C features of the form “feature= c ?” for every $c \in C$.
↑ binary features!

Q: How about features of the form “feature $\in C'$ ” for every $C' \subset C$?

computational complexity ↑
Because there are 2^C subsets!

- Real value e.g., weight, age
 - Sort the values.
 - Find the **breakpoints**: For every two adjacent points with opposite labels, compute the midpoint.
 - Derive features like “weight \leq breakpoint”



- Binary
- Multiclass: What changes do we need to make?
 - Almost none! Just extend the computation of accuracy to multiclass.

If the number of classes is >2

Notions of uncertainty: general case

Suppose in $S \subseteq \mathcal{X} \times \mathcal{Y}$, a p_k fraction are labeled as k (for each $k \in \mathcal{Y}$).

① Classification error:

$$u(S) := 1 - \max_{k \in \mathcal{Y}} p_k$$

② Gini index:

$$u(S) := 1 - \sum_{k \in \mathcal{Y}} p_k^2$$

③ Entropy:

$$u(S) := \sum_{k \in \mathcal{Y}} p_k \log \frac{1}{p_k}$$

Each is *maximized* when $p_k = 1/|\mathcal{Y}|$ for all $k \in \mathcal{Y}$
(i.e., equal numbers of each label in S)

Each is *minimized* when $p_k = 1$ for a single label $k \in \mathcal{Y}$
(so S is **pure** in label)

Regression: when the labels are real numbers

64

- Classification vs Regression
 - Both supervised learning
 - Regression has real-valued labels.
- Examples: Price prediction. Property value prediction.
- Standard measure of performance: mean squared error: $\frac{1}{m} \sum_{i=1}^m (f(x^{(i)}) - y^{(i)})^2$

Q: why are we using squared error rather than absolute error?

my opinion: convenience & tradition

- What changes needed for decision tree?

- How to make predictions at the leaf node?

Average labels of the data at the leaf;
denote by \bar{y}_{YES} and \bar{y}_{NO} .

- How to adjust score[f]?

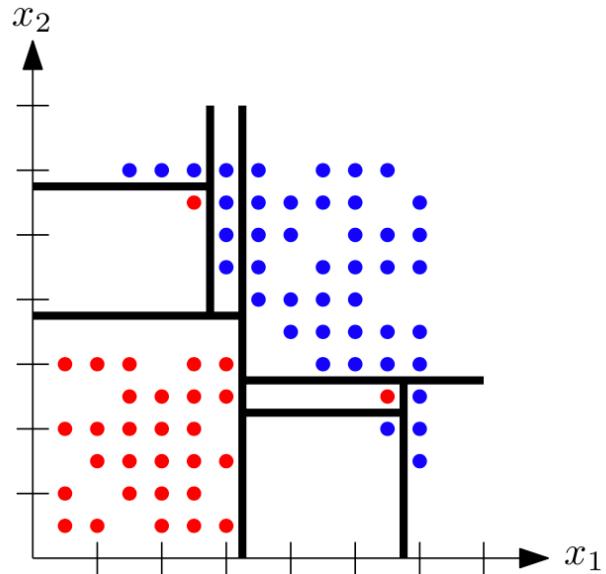
Use negative squared error

$$\frac{1}{|data|} \cdot \left(- \sum_{i \in YES} (y_i - \bar{y}_{YES})^2 - \sum_{i \in NO} (y_i - \bar{y}_{NO})^2 \right)$$

(notations from the decision tree pseudocode)

“Spurious” patterns can be learned

66



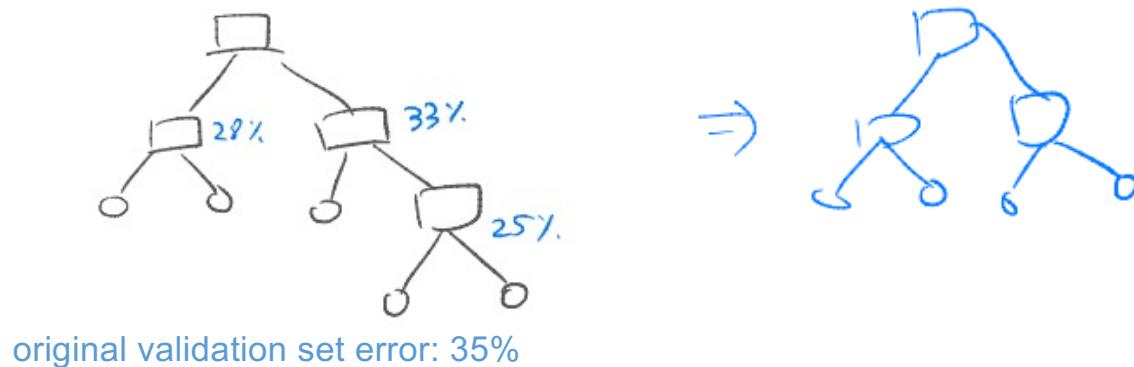
by the way, note axis-parallel decision boundaries

Unlearn spurious patterns by pruning

67

Split the data into **train set** and **validation set**

- Build a decision tree based on the **train set**; compute the **validation set** error
- While true
 - For each non-leaf node, pretend that it is a leaf node and then compute the validation set error (but do not make it a leaf node yet)
 - If none reduces the validation set error
 - Break
 - Else
 - **Prune** the one that reduces the validation set error the most

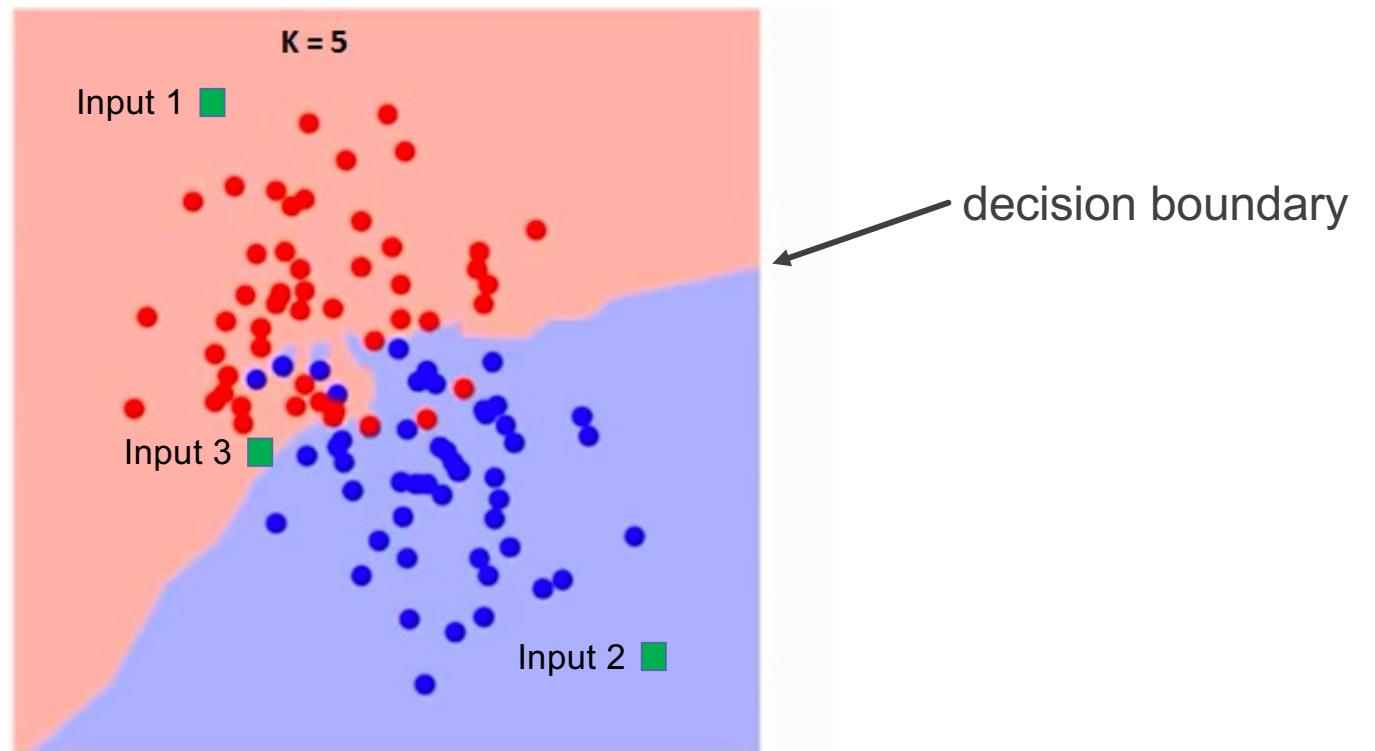


k-Nearest Neighbors (k-NN)

- Train set: $S = \{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$
- **Idea:** given a new, unseen data point x , its label should resemble the labels of **nearby points**
- What function?
 - Input: $x \in \mathbb{R}^d$
 - From S , find the k nearest points to x from S ; call it $N(x)$
 - E.g., Euclidean distance
 - Output: the majority vote of $\{y_i : i \in N(x)\}$
 - For regression, take the average label.

k-NN example

70



How to extract features as real values?

- Binary features: Take 0/1
- Categorical $\{1, \dots, C\}$ (e.g., movie genres)
 - Binary vector of length C . Set c -th coordinate 1 and 0 otherwise. one-hot encoding

Q: Why don't we just take $1, \dots, C$ as a real-valued feature?

Distance:

- (popular) Euclidean distance: $d(x, x') = \sqrt{\sum_{i=1}^d (x_i - x'_i)^2}$
- Manhattan distance : $d(x, x') = \sum_{i=1}^d |x_i - x'_i|$

Q: If we shift a feature, would the distance change? no

Q: What about scaling a feature? yes

- Features having different scale can be problematic. (e.g., weights in lbs vs shoe size)
- [Definition] **Standardization**

- For each feature f , compute $\mu_f = \frac{1}{m} \sum_{i=1}^m x_f^{(i)}$, $\sigma_f = \sqrt{\frac{1}{m} \sum_{i=1}^m (x_f^{(i)} - \mu_f)^2}$
- Then, transform the data by $\forall f \in \{1, \dots, d\}, \forall i \in \{1, \dots, m\}, x_f^{(i)} \leftarrow \frac{x_f^{(i)} - \mu_f}{\sigma_f}$

after transformation, each feature has mean 0 and variance 1

- Be sure to keep the “standardize” function and apply it to the test points.
 - Save $\{(\mu_f, \sigma_f)\}_{f=1}^d$
 - For test point x^* , apply $x_f^* \leftarrow \frac{x_f^* - \mu_f}{\sigma_f}, \forall f$

- Given: labeled data D
- Training
 - Compute and save $\{(\mu_f, \sigma_f)\}_{f=1}^d$
 - Compute and save standardization of D
- Test
 - Given x^* , apply standardization $x_f^* \leftarrow \frac{x_f^* - \mu_f}{\sigma_f}, \forall f$
 - Compute k nearest neighbors $N(x^*)$
 - Predict by majority vote label in $N(x^*)$ (average label for regression tasks)

Variations

74

Recall the majority vote rule: $\hat{y} = \arg \max_{y \in \{1, \dots, C\}} \sum_{i \in \mathcal{N}(x)} 1\{y_i = y\}$

⊕
⊕

○
⊖

Q: Blue dot is the test point. If k=3, which label would it predict?

Q: Which label do you think we should predict?

Weighted version

- $\hat{y} = \arg \max_{y \in \{1, \dots, C\}} \sum_{i \in \mathcal{N}(x)} w_i 1\{y^{(i)} = y\}$

weights that sum to 1

$$w_i \propto \exp(-\beta \cdot d(x, x^{(i)})), \beta > 0$$

~~$$w_i \propto \frac{1}{d(x, x^{(i)})^\beta}$$~~

$$w_i \propto \frac{1}{1 + d(x, x^{(i)})^\beta}$$

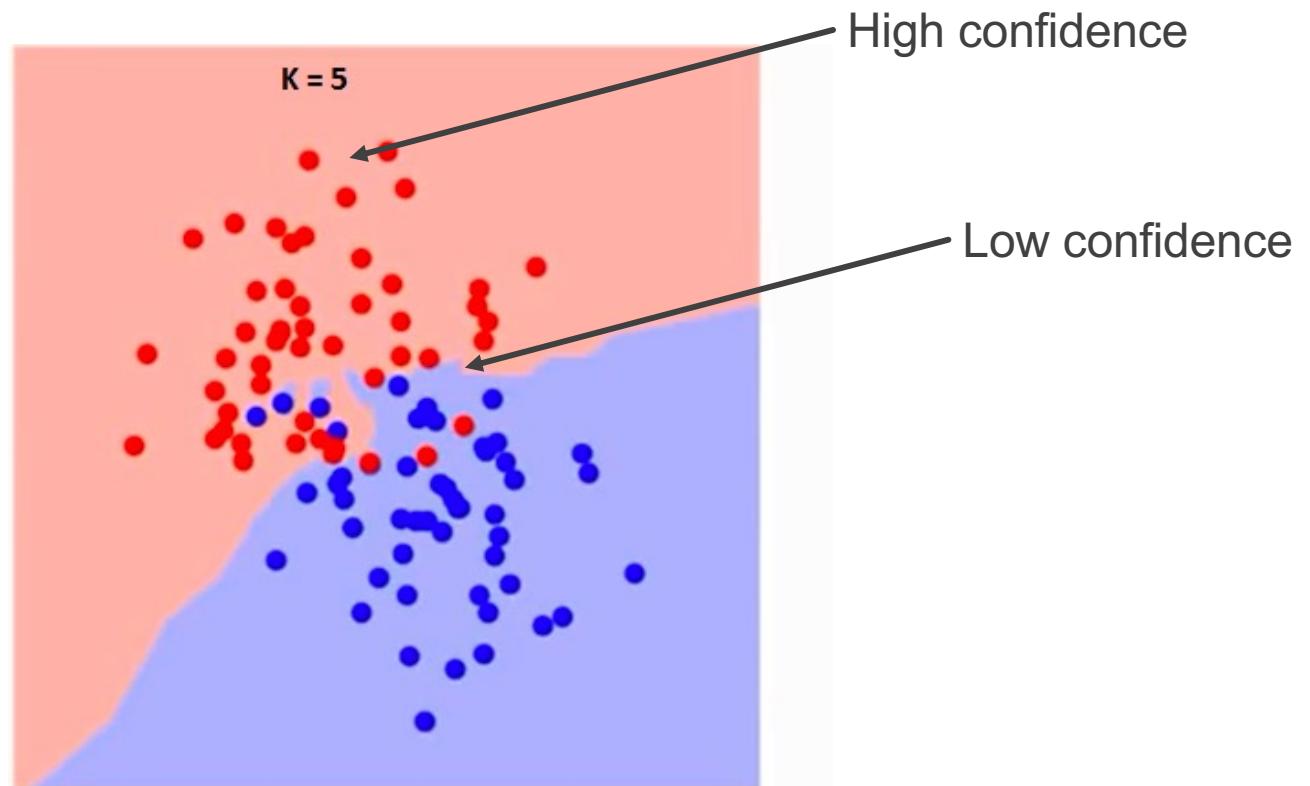
Q: What would be the downside of using weighted version?

tuning β is cumbersome!

Confidence

- $P(Y = y|X = x) \propto \sum_{i \in \mathcal{N}(x)} 1\{y^{(i)} = y\}$
- $P(Y = y|X = x) \propto \sum_{i \in \mathcal{N}(x)} w_i 1\{y^{(i)} = y\}$ // weighted version

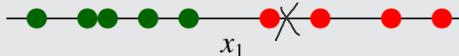
Same thing applies to decision tree – (number of majority points in that leaf node / number of points in that leaf node)



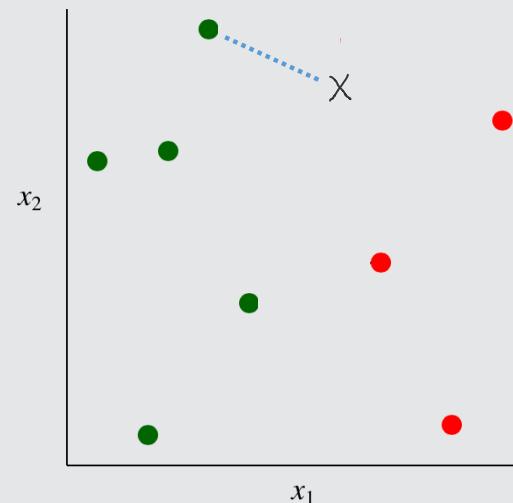
Issues 1: irrelevant features

77

here's a case in which there is one relevant feature x_1 and a 1-NN rule classifies each instance correctly



consider the effect of an irrelevant feature x_2 on distances and nearest neighbors



Q: how did we deal with irrelevant features in decision trees?

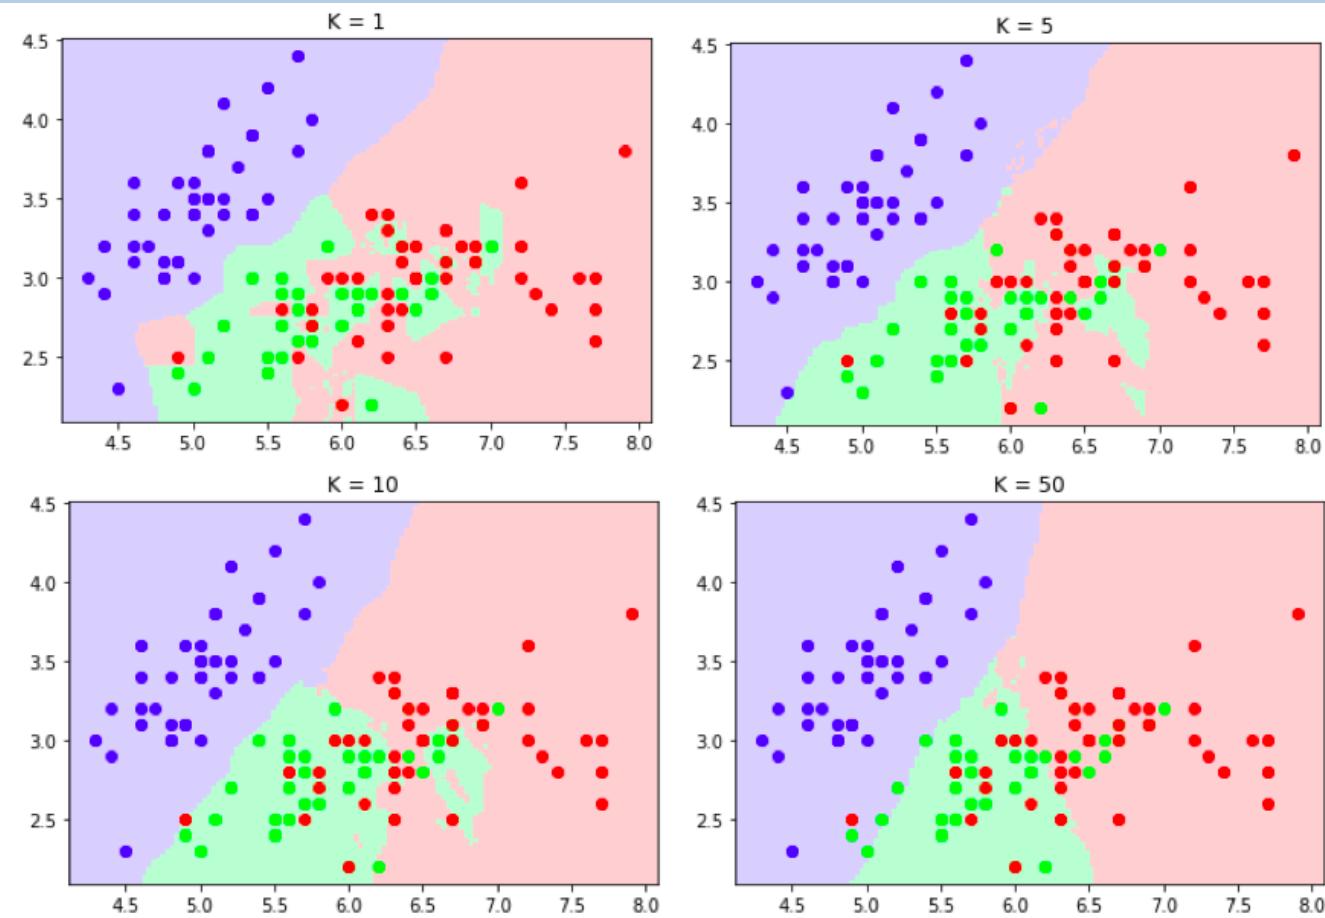
not all features are used because

- (i) we stop adding features when they are unnecessary (e.g. having zero local accuracies, subset is already pure)
- (ii) pruning

- How a k-NN function work:
 - Compute distance to m points $O(dm)$
 - Sort distances $O(m \log m)$
 - Pick k smallest. $O(k)$
 - Overall $O(m(d + \log m))$
- Issue: test time complexity scales linearly with m !!
- Solutions
 - k-d tree: Exact search
 - Best case: $O(\log(m))$ Worst case: $O(m)$

for large d very likely to hit the worst case
 - Locality-sensitive hashing: approximate search, $O(m^\rho)$ with $\rho \in (0,1)$

- Q: If we set $k = m$, then which classification rule does it look like?
- Q: If we set $k = 1$, what would be the train set error (assume there is no repeated train data point)?



Comparison

81

	Decision Tree	k-NN
• Interpretability	good	bad
• Sensitivity to irrelevant features	low	high
• train time	$O(dm^2 + dm \log m)$	$O(dm)$
• test time	depth of the tree worst: $O(\min\{d, m\})$ best: $\log(m)$	$O(m(d + \log(m)))$ bad
• test time space complexity	worst: $O(m)$ in general: much smaller	$\Theta(dm)$

- Model selection
 - How to choose k?
 - Overfitting
- Naïve Bayes Classifier

Thank you!