



Computer  
Science

# **CSC380: Principles of Data Science**

**Advanced ML algorithms**

Kyoungseok Jang

## Announcement

- Reading quiz today
  - Already announced quiz questions on Piazza.
- Uploaded Final practice exam.

## Outline

- Random Forest (supervised learning, advanced decision tree)
- Recommendation algorithms (unsupervised learning)
- Generative Adversarial Network (generative model)

## Disclaimer

- Today's lecture will not be included in your final exam.
- You don't need to study this part too hard.

# Random Forest

# Review: Decision Tree

---

## Algorithm 1 DECISIONTREETRAIN(*data*, *remaining features*)

---

```


1: guess ← most frequent answer in data           // default answer for this data
2: if the labels in data are unambiguous then           <= i.e., all data points have the same
3:   return LEAF(guess)                                   // base case: no need to split further           label
4: else if remaining features is empty then
5:   return LEAF(guess)                                   // base case: cannot split further
6: else                                                   // we need to query more features
7:   for all  $f \in \text{remaining features}$  do           <= there is no point in adding a
8:     NO ← the subset of data on which  $f=no$            feature that appeared in its parent!
9:     YES ← the subset of data on which  $f=yes$ 
10:     $score[f] \leftarrow ( \# \text{ of majority vote answers in NO} \quad <= \text{answer} = \text{label}$ 
11:      -  $+ \# \text{ of majority vote answers in YES} ) /$ 
12:       $size(data)$ 
13:   end for
14:   f ← the feature with maximal  $score(f)$ 
15:   NO ← the subset of data on which  $f=no$ 
16:   YES ← the subset of data on which  $f=yes$ 
17:   left ← DECISIONTREETRAIN(NO, remaining features \ {f})
18:   right ← DECISIONTREETRAIN(YES, remaining features \ {f})
19:   return NODE(f, left, right)

```

---


# Advantages of Decision Tree

- Pros
  - Interpretability
  - Less data preparation (preprocessing)
  - Non-parametric: not like Naïve-Bayes, it does not require complicated model assumption.
  - Versatility
  - Non-linearity



# Disadvantages of Decision Tree

- Cons
  - **Overfitting: model can be complex = vulnerable to overfitting**
  - **Optimization: order  $O(dm^2 + dm \log m)$** 
    - especially for the continuous feature → Too many features to consider.

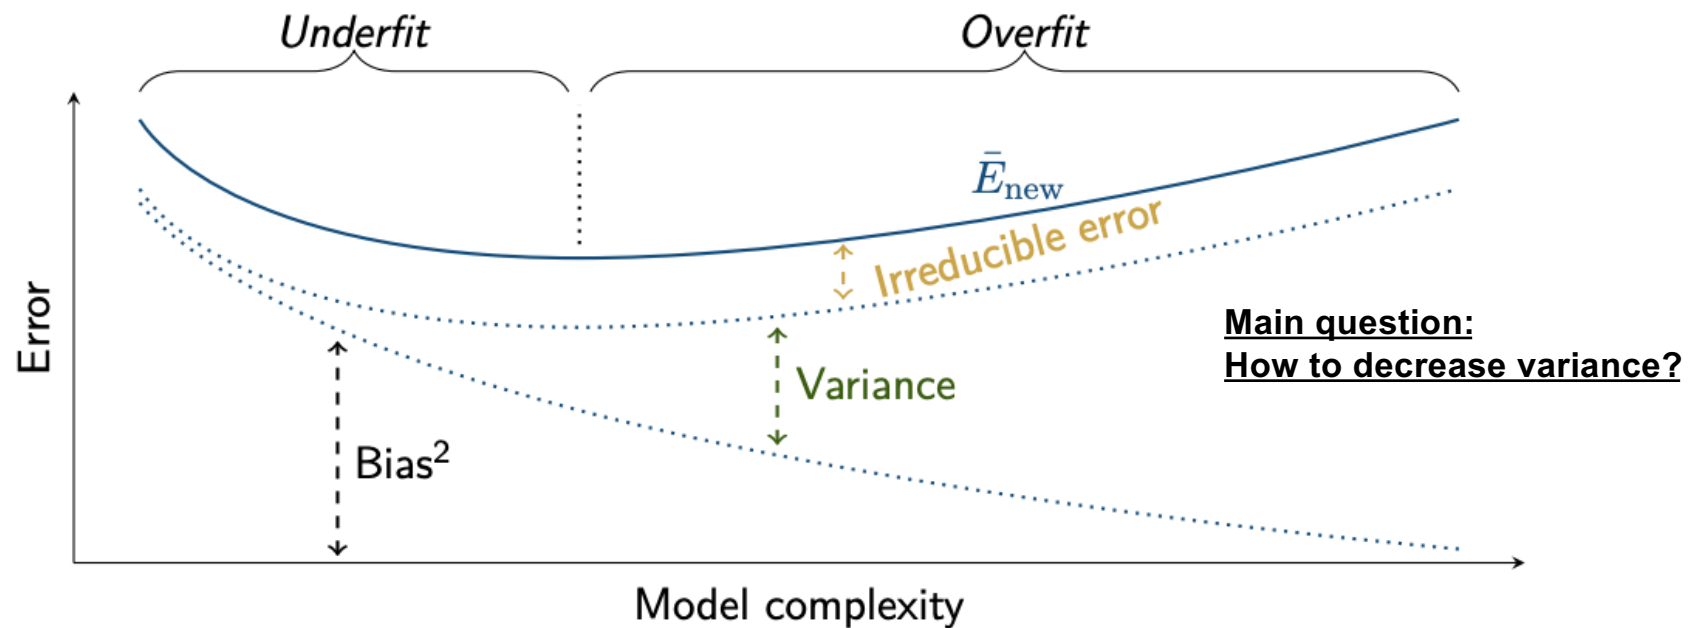


Pruning is usually not enough  
+ Additional complicated computation





# Bias-Variance Tradeoff



Finding a balanced fit (neither over- nor underfit) is called the **the bias-variance tradeoff**.

## Bias-Variance Tradeoff

- Intuition: For i.i.d random variables  $X_1, \dots, X_n$

$$\text{Var}\left(\frac{1}{n} \sum_{i=1}^n X_i\right) = \text{Var}(X_1)$$

- Observation: If we have D independent (and large enough) dataset, then we can train D individual decision trees and do 'majority vote' to decrease variance!
- Problem
  - D datasets should be independent! We usually don't have enough data to split it by D large enough subsets!
  - Alternative: Bootstrapping

# Bagging

- Bagging = Bootstrap + Aggregating
- **Objective: Create multiple decision trees**

- **Bootstrapping:**

Generate multiple samples of training data via bootstrapping (sampling from your dataset)


Each tree will be trained based on this sampled dataset.

- **Aggregating:** For a given input, we output
  - Regression: the averaged outputs of all the models for that input.
  - Classification: the class that is outputted by the majority

training examples						
	#1	#2	#3	#4	#5	#6
original dataset	1	1	1	1	1	1
decision tree 1	1	1	0	2	1	1
decision tree 2	3	0	1	0	2	0
decision tree 3	0	1	3	1	0	1

# Bagging

- This is one example of **ensemble method**
  - Method of building a single model by training and aggregating multiple models



## Out-of-bag evaluation

- How to evaluate generalization errors?
  - Validation set? Yes, it works, but... it would be better if we don't split!
  - Turns out we don't need to split!
  - Traditionally, ~67% of original data is in a sampled dataset.
  - You can use the rest of the data as your validation set for each tree!

	Training examples						Examples for OOB Evaluation
	#1	#2	#3	#4	#5	#6	
original dataset	1	1	1	1	1	1	
decision tree 1	1	1	0	2	1	1	#3
decision tree 2	3	0	1	0	2	0	#2, #4, and #6
decision tree 3	0	1	3	1	0	1	#1 and #5


## Good, are we done?

- Not yet... are those trees 'truly' independent?
- These ensemble models are meaningless when models are dependent on each other.
  - E.g.)  $X_1, \dots, X_n$ : All equal random variable ( $X_i = X$ )
  - Then  $\text{Var}(\frac{1}{n} \sum_{i=1}^n X_i) = \text{Var}(X) = \text{Var}(X_1)$ : variance does not decrease!
- In practice, these ensembles of trees in Bagging tend to be highly correlated!
  - E.g.) One extremely strong predictor,  $x_j$ , in the training set amongst moderate predictors.
  - Then most of your trees will start with  $x_j$ , and no big changes.
  - Multiple identical trees!

# Random Forest

- How do make trees 'different' from each other?
  - Don't allow your algorithm to use all the features!
- 1. Create separate bootstrap samples (same as bagging)
- 2. For each tree, at each split (node training), **we randomly select a set of  $J'$  predictors from the full set of predictors.**
- 3. Find the best feature among  $J'$

Blue: randomly selected candidate  $J'$   
Red box: chosen feature for this node



## Random Forest - properties

- Parallel training is possible → Not that heavier optimization than one decision tree
- Three main hyperparameters to tune
  - Number of predictors to randomly select for each split (node)
  - Depth of the tree, or the minimum leaf node size (complexity of tree)
  - Number of trees: turns out, increasing number of trees **does not increase variance!**
- When the number of predictors is large, but the number of relevant predictors is small, random forests can perform poorly.



## Pros and Cons

- Pros

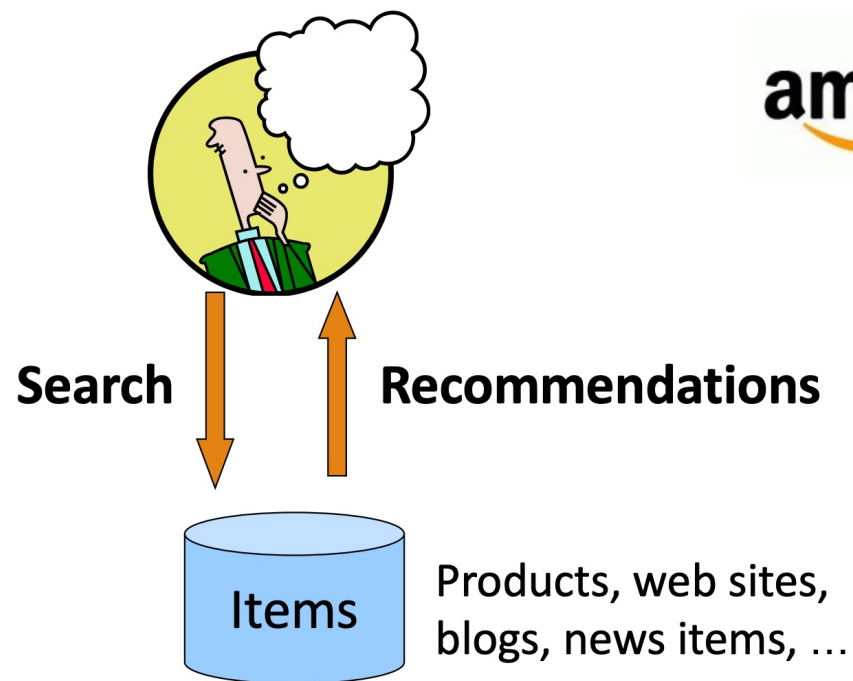
- **Hard to overfit:** "[Random Forests do not overfit](#)" – Leo Breiman
- Even without worrying about overfitting, one can create a complex model!
- Inherits most of the advantages of decision trees
  - Less data preparation (preprocessing)
  - Non-parametric
  - Versatility
  - Non-linearity
- Can be trained parallel - Faster learning relative to its size

- Cons


- **Lose Interpretability.**
- **The size of the model can be large.** Models with more than 1M nodes are common.
- Random forests cannot learn and reuse internal representations.

# Recommendation algorithms

# Recommendations are everywhere



## Types of recommendations

- Editorial and hand-curated
  - List of 'essential' by XYZ
  - Personal recommendation of critic ABC
- Simple Aggregation
  - Top 100, recent uploads, most popular
- Tailored to individual users
  - Amazon, Youtube, Netflix, etc...  Today's main topic

## The task

- Given the information of a user, what should we recommend next?
- Formal model
- $X$ =set of users
- $S$ =set of items
- Utility function  $u: X \times S \rightarrow R$ 
  - $R$ : set of ratings (e.g. 1-5 stars, real numbers in  $[0,1]$ )
- Want to maximize customer's happiness! (so that the company can earn more)

## The task

- This model naturally induces ‘utility matrix’

		Harry Potter			Twilight		Star Wars	
		HP1	HP2	HP3	TW	SW1	SW2	SW3
Anita	<i>A</i>	4			5	1		
Beyonce	<i>B</i>	5	5	4				
Calvin	<i>C</i>				2	4	5	
David	<i>D</i>		3		?			3

- How to fill out the ‘empty spaces?’
- If we can ‘guess’ the numbers in empty spaces, we can recommend best item for the customer?

## Two approaches

- Content-Based Filtering
  - Recommend items to customer  $x$  similar to previous items rated highly by  $x$ .
  - User-Item interaction
- Collaborative Filtering
  - Find a set of other users whose ratings are “similar” to  $x$ ’s ratings
  - 1) User-user interaction (straightforward)
  - 2) Item-item interaction
  - 3) Model-based: google developer version

## Content-based filtering

- For each item, create an item profile (vector)
  - E.g.) Movie: genre, director, actor, year, ...

	Melissa McCarthy	Actor A	Actor B	...	Johnny Depp	Comic Genre	Spy Genre	Pirate Genre
Movie X	0	1	1	0	1	1	0	1
Movie Y	1	1	0	1	0	1	1	0





- For each user, create a user profile (vector) with same features!
  - E.g.) Average out all movies that user have seen.

	Melissa McCarthy	Actor A	Actor B	...
User U	0.2	.0005	0	0



## Content-based filtering

- Measure the ‘similarity score’
  - Various types of scores: dot product, cosine similarity, Pearson similarity, Euclidean distance
- Recommend an item with the highest similarity score

		Education	Casual	Health		TimeWastr	Science R Us	Healthcare
Item 1		●					●	
Item 2			●		...	●		
Item 3				●				●
User		●			...		●	●

If we use the dot product as our similarity score, what item should we recommend to the user?

## Pros and Cons

- Pros
  - No need for data from other users
  - Able to recommend to users with unique tastes
    - No first-rater problem: if the movie is well featured
  - Able to provide an explanation: 'Since you liked ...'
- Cons
  - How to determine features appropriately?
  - Frequently involves hand-engineering
  - **Overspecialization** – hard to introduce or expand new content
    - Never recommend content outside the user's profile

## Collaborative Filtering

- Instead of somewhat hand-engineered features, we will use the utility matrix directly!
  - Since it uses a real dataset, it will reflect reality better.
- But how?
  - 1) User-to-user interaction: directly compute using rows!
  - 2) Item-to-item interaction: directly compute using columns!
  - 3) Model-based interaction: create a content-based filtering model latently!


		Harry Potter			Twilight		Star Wars	
		HP1	HP2	HP3	TW	SW1	SW2	SW3
Anita	A	4			5	1		
Beyonce	B	5	5	4				
Calvin	C				2	4	5	
David	D		3					3

## Collaborative Filtering – User-user based


- Basic concept: Find similar users and recommend items that they like!
- How to define similarity?
  - Use the  $i$ -th row of the utility table (items that user  $i$  has used) as the feature vector of user  $i$
  - Use the similarity measure we explained before!
    - Cosine measure, dot product, Pearson similarity, ...
    - There are lots of details added for this process, but I will skip them.
- The importance of the user is proportional to the similarity.
  - Because it is highly probable that similar people with input user  $x$  will tell you more about the taste of  $x$ .

## Collaborative Filtering – User-to-user based

- There are lots of variants for this user-user based CF, but here's one example:
- $r_x$ : the vector of the user x's ratings
- $N$ : be the set of k 'neighborhoods' (k most similar to x)
- Prediction for the item i of user x (can be vary):
- $$r_{xi} = \frac{\sum_{n \in N} \text{sim}(x,n) r_n}{\sum_{n \in N} \text{sim}(x,n)}$$
- Recommended item  $i^* = \arg \max_i r_{xi}$



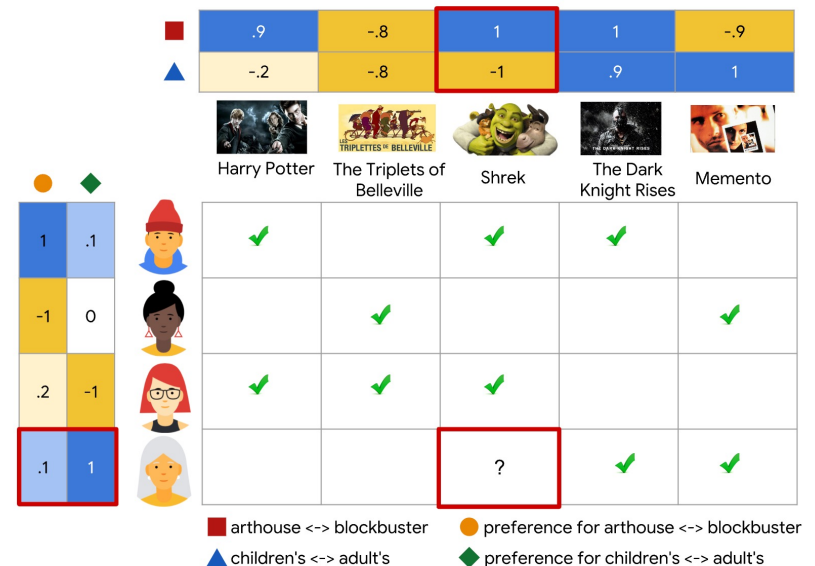
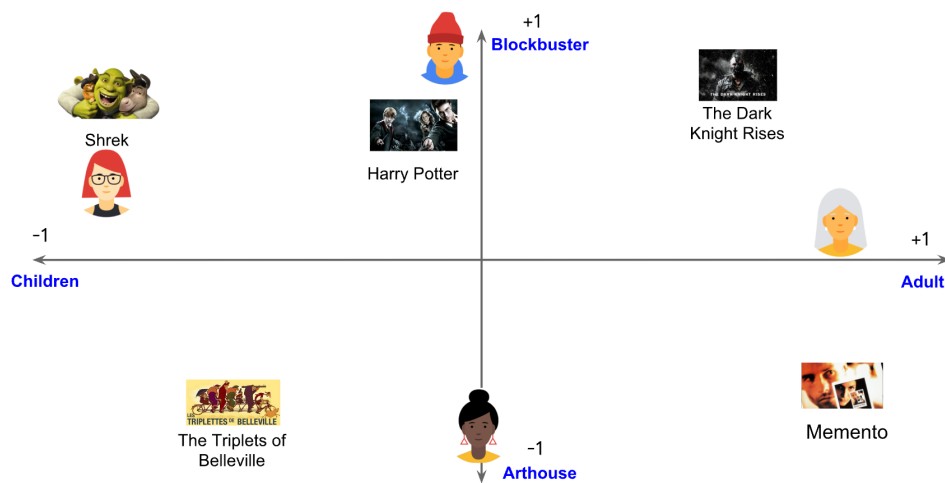
**(a) User-based filtering**



**(b) Item-based filtering**

# Collaborative Filtering: Model-based

- Imagine that we already have hidden 'embedding' – feature vectors that affects utility.
- Then our prediction will be easy – as we did in Content-based filtering. E.g.) Inner product between user vec and item vec!



# Collaborative Filtering: Model-based

- It will be great if we can determine the features ‘automatically’
- But how? Using matrix factorization...
- Suppose that there exists a matrix  $U \in \mathbb{R}^{m \times d}$  (features of users) and  $V \in \mathbb{R}^{n \times d}$  (features of items) such that  $UV^T$  represents the true utility matrix.





## Collaborative Filtering: Model-based

- If this model is true, then it should satisfy
  - If the actual rating of user  $i$  for item  $j$  is  $A_{ij}$ , then this model should predict similar rating for this user-item pair, which means
    - $(A_{ij} - \langle U_i, V_j \rangle)^2$  should be small for all observed  $i$  and  $j$ .
  - For unobserved items, it might 'implicitly' mean that the user is not interested in that items.
    - So, it would be great  $(0 - \langle U_i, V_j \rangle)^2$  is also small for unobserved

- You need to optimize

$$\arg \min_{U \in \mathbb{R}^{m \times d}, V \in \mathbb{R}^{n \times d}} ||A - UV^T||_F^2$$

Where  $||B||_F^2 = \sum b_{ij}^2$ , sum of squares of all entries in matrix.

- Turns out, it can be done efficiently!

## Pros and Cons

- Pros

- **No domain knowledge:** everything can be done automatically and mathematically. Only needs the feedback matrix!
- **Serendipity:** The model will help you discover your new interests based on the interest of other users.
- Relatively robust and performs well in practice.

- Cons

- **Cold-start problem:** new user or new item – no similarity to use!
- **Interpretability:** all the features are created ‘automatically’, so it is hard to understand what each feature means.
- **Popularity bias:** tend to recommend popular one – someone with unique taste will not like it.
- Sometimes ethical issues...